

Explainable Artificial Intelligence

03 – Neuro-symbolic approaches

MSc in Artificial Intelligence

MSc in Computer Engineering

Marco Lippi

marco.lippi@unifi.it



UNIVERSITÀ
DEGLI STUDI
FIRENZE

Table of Contents

1 General concepts

- ▶ General concepts
- ▶ Combining learning and reasoning
- ▶ Tasks
- ▶ Models and Frameworks

The connectionist vs. symbolic dilemma

1 General concepts

A central question in AI: how is knowledge represented in our mind?

Symbolic approaches

- Reasoning as the result of formal manipulation of symbols

Connectionist (sub-symbolic) approaches


- Reasoning as the result of processing of interconnected (networks of) simple units

Connectionist vs. symbolic AI

1 General concepts

Symbolic approaches

- Founded on the **principles of logic**
- Exploiting **background knowledge**
- Highly **interpretable**
- Need **discretization**



toxic(M) :- doublebond(M,C1,C2), hydroxyl(C2), methyl(M)

Connectionist vs. symbolic AI

1 General concepts

Connectionist approaches

- Can more easily deal with **uncertain knowledge**
- No need to **encode** knowledge at all!
- Need to observe (large) **data** collections!
- Algorithms typically can be easily **distributed**
- Often seen as “black box” → dark magic
- **Graceful degradation** of performance with noisy input

Connectionist vs. symbolic AI

1 General concepts

Both a technological and a philosophical difference...

- Do we (as humans) really need huge data collections to learn?
- How do we acquire **generalization** skills?
- Can we easily define **rules** that drive our decisions?
- Can we easily describe data through a symbolic formalism?


Table of Contents

2 Combining learning and reasoning

- ▶ General concepts
- ▶ Combining learning and reasoning
- ▶ Tasks
- ▶ Models and Frameworks

Combining learning and reasoning

2 Combining learning and reasoning




[Image by Luc De Raedt]

Combining learning and reasoning

2 Combining learning and reasoning


Machine (and deep) learning



Combining learning and reasoning

2 Combining learning and reasoning

Statistical Relational Learning (SRL/StarAI)




[Image by Luc De Raedt]

Combining learning and reasoning

2 Combining learning and reasoning

Neural Symbolic Computation (NeSy)




[Image by Luc De Raedt]

Pioneering approaches: KBANNs

2 Combining learning and reasoning

Knowledge-based artificial neural networks [Towell & Shavlik, 1994]

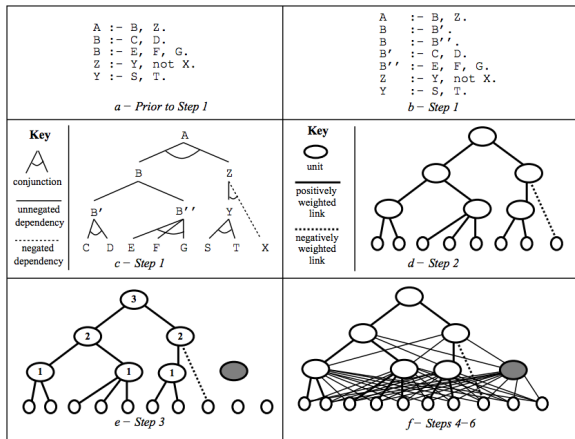
- One of the first attempts to inject knowledge into ANNs
- Trying to interpret an ANN model as logic rules



<i>Knowledge Base</i>		<i>Neural Network</i>
Final Conclusions	\longleftrightarrow	Output Units
Supporting Facts	\longleftrightarrow	Input Units
Intermediate Conclusions	\longleftrightarrow	Hidden Units
Dependencies	\longleftrightarrow	Weighted Connections

Pioneering approaches: KBANNs

2 Combining learning and reasoning



NeSy and SRL/StarAI

2 Combining learning and reasoning

More recent research directions:

- **Statistical Relational Learning (SRL) or Statistical Relational AI (StarAI)**
- **Neural-Symbolic Learning and Reasoning (NeSy)**

Both developed during the 90s-00s

The former combining logic with probabilistic/statistical learning (SRL/StarAI)

The latter combining logic with cognitive neuroscience (NeSy)

Which logic?

Which probabilistic models?

Which neural models?

Aiming to combine **first-order logic** and **graphical models** for learning and reasoning

- Exploit the **expressive power** of first-order logic
- Handle uncertainty with **graphical models**
- Combine **logic and probabilistic** inference

Aiming to combine **neural models** and **symbolic approaches** for learning and reasoning

- Encode knowledge in the **architecture** of the network
- Use a **regularization** term to encode rules
- **Constrain** neural computations with rules

Caveat: injecting knowledge into the neural network, then let the network do the rest might not be sufficient (partly lost the power of reasoning and explanation)

Logic can be used in several ways [De Raedt et al., 2020]:

- Logic as a kind of neural program
- Logic as a regularizer
- ...

Neuro-Symbolic AI

2 Combining learning and reasoning

Logic as a kind of neural program

Logic cabled within the architecture of the network

This comes from the idea of KBANNs by Towell and Shavlik...

Logic as **regularizer**

Combine standard classification loss with a loss that takes into account constraints, by penalizing solutions that break them...


$$\text{Loss} = \text{ClassificationLoss} + \lambda \cdot \text{SemanticLoss}$$

But... If logic is encoded in the network, how to **reason logically**?

Neuro-Symbolic AI

2 Combining learning and reasoning

multi-class classification



This constraint should be satisfied

$$(\neg x_1 \wedge \neg x_2 \wedge x_3) \vee$$

$$(\neg x_1 \wedge x_2 \wedge \neg x_3) \vee$$

$$(x_1 \wedge \neg x_2 \wedge \neg x_3)$$

from Xu et al., ICML 2018

Neuro-Symbolic AI

2 Combining learning and reasoning

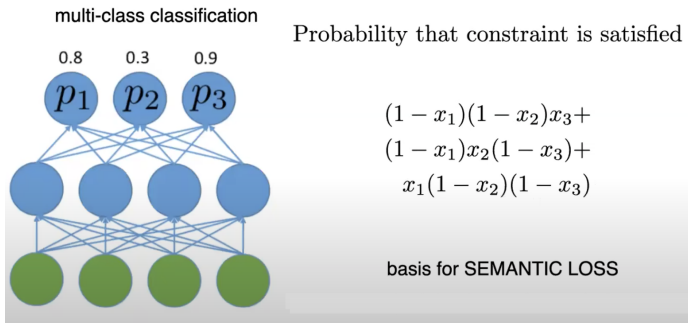


Table of Contents

3 Tasks

- ▶ General concepts
- ▶ Combining learning and reasoning
- ▶ **Tasks**
- ▶ Models and Frameworks

Classic SRL/StarAI tasks

3 Tasks

StarAI applications typically deal with three distinct, but strongly inter-related problems...

- Inference
- Parameter Learning
- Structure Learning

Inference

3 Tasks

Inference in StarAI lies at the intersection between logical and probabilistic inference

Logical Inference

Inferring the truth value of some logic facts, given a collection of other facts and rules

Probabilistic inference

Inferring the posterior distribution of unobserved random variables, given observed ones

Logical theory

```
likesmovie(X,M) :- moviegenre(M,G), likesgenre(X,G).  
likesmovie(X,M) :- friends(X,Y), likesmovie(Y,M).
```

Caveat: need of grounding out...

```
likesmovie(alice,bladerunner).  
friends(alice,bob).  
likesgenre(alice,sciencefiction).  
...
```

Main techniques

- Model-based, SAT solvers
- Proof-based, finding derivations

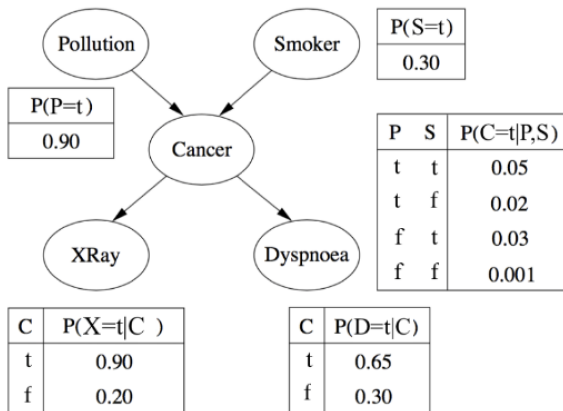
Grounding is in any case the **bottleneck**...

- Smart techniques to avoid generating **all** the predicates
- Exploit symmetries and patterns/templates

Probabilistic inference

3 Tasks

An example in Bayesian Networks



Parameter learning

3 Tasks

Typically, StarAI models specify a set of **parameters** (probabilities or real values) attached to rules/clauses to model uncertainty in the available knowledge

These parameters can be **learned from data**:

- probability tables in Bayesian networks
- weights or probabilities attached to soft rules

Structure learning

3 Tasks

Highly challenging problem: directly learning the rules (the **structure**) of the model

Different approaches...

- Jointly learn parameters and rules
- First learn rules (i.e., with ILP), then their weights

Collective classification

3 Tasks

This framework could be easily exploited to perform collective classification on a set of **non-independent examples**, like nodes in a graph or sentences in a document, ...

Given a set of (possibly neural) rules, and a collection of constants/features representing the document, the inference algorithm computes the **most likely world**, or interpretation, thus assigning a truth value to each predicate in the document.




Table of Contents

4 Models and Frameworks

- ▶ General concepts
- ▶ Combining learning and reasoning
- ▶ Tasks
- ▶ **Models and Frameworks**

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

Logic imposes **hard constraints** on the set of possible worlds

Markov logic exploits **soft constraints**

A Markov Logic Network is defined by:

- a set of first-order formulae
- a set of weights, one attached to each formula

A world violating a formula is less probable, **not** impossible!

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

A probabilistic-logic framework to model knowledge

Syntax **opposite to Prolog**: constants uppercase, variables lowercase

An example

```
Movie = {BladeRunner, TheMatrix}
```

```
Person = {Alice, Bob, Carl, David}
```

```
2.3 LikesMovie(x,m)  $\wedge$  Friends(x,y)  $\Rightarrow$  LikesMovie(y,m)
```

```
1.6 Friends(x,y)  $\wedge$  Friends(y,z)  $\Rightarrow$  Friends(x,z)
```

The **higher** the weight, the **more likely** a world where rule is true, other things being equal

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

Beware of the differences in the syntax:

- MLN: uppercase constants (e.g., Alice) and lowercase variables (e.g., person)
- ProbLog: lowercase constants (e.g., alice) and uppercase variables (e.g., Person)

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

The semantics of MLNs induces a **probability distribution** over all possible worlds. We indicate with X a set of random variables represented in the model, then we have:

$$P(X = x) = \frac{\exp(\sum_{F_i \in \mathcal{F}} w_i n_i(x))}{Z}$$

being $n_i(x)$ the number of true groundings of F_i in world x and Z is the partition function:

$$Z = \sum_{x \in \mathcal{X}} \exp \left(\sum_{F_i \in \mathcal{F}} w_i n_i(x) \right)$$

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

Discriminative setting

Typically, some atoms are always **observed** (evidence X), while others are **unknown** at prediction time (query Y)

EVIDENCE	QUERY
Friends(Alice,Bob)	LikedMovie(Alice,TheMatrix)???
Friends(Alice,Carl)	LikedMovie(Alice,PulpFiction)???
WatchedMovie(Alice,PulpFiction)	LikedMovie(Bob,TheMatrix)???
WatchedMovie(Bob,PulpFiction)	LikedMovie(Bob,PulpFiction)???

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

The **probability** of a world/configuration depends on the **weights** (w_i) and the number of **groundings** (n_i) of each formula (F_i):

$$P(Y = y|X = x) = \frac{\exp(\sum_{F_i \in \mathcal{F}} w_i n_i(x, y))}{Z_x}$$

Inference aims to find the **most probable** y given x :

$$y^* = \operatorname{argmax}_y P(Y = y|X = x)$$

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

Inference

Given set of known facts, weighted rules used to infer truth value of other (query) facts

```
LikesMovie(Alice,BladeRunner)
Friends(Alice,Bob)
Friends(Alice,Carl)
```

```
LikesMovie(Carl,BladeRunner)???
```

#P-complete problem \Rightarrow approximate algorithms

MaxWalkSAT [1996], **stochastic local search** \Rightarrow minimize the sum of unsatisfied clauses

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

A common problem in (probabilistic) logical inference is that you typically need to **ground** your knowledge base into your predicates.

This can quickly become **unaffordable** due to both time and memory requirements...

...Now moving towards **lifted inference**!

Markov Logic Networks [Richardson & Domingos, 2006]

4 Models and Frameworks

Learning

Both **weights** and **rules** themselves can be **learned** from a collection of predicate observations.

Maximize conditional log likelihood (CLL) of query predicates given evidence: **inference as subroutine!**

$$\frac{\partial}{\partial w_i} \log P(Y = y | X = x) = n_i - E_w[n_i]$$

Ground-Specific MLNs [Lippi & Frasconi, 2009]

4 Models and Frameworks

An extension of MLNs that allows to embed neural networks to compute weights

A simple classification example

$w(x)$ `HasFeatures(x,$f) => PositiveClass(x)`

The weight $w(x)$ is computed by a neural network using (any) set of features f describing example x .

These are named **Ground-Specific MLNs**.

Ground-Specific MLNs [Lippi & Frasconi, 2009]

4 Models and Frameworks

An example in structured text classification

```
2.3  Features(X,$F1) => CategoryA(X)
-1.8  Features(X,$F1) => CategoryB(X)
0.9   Features(Y,$F2) => CategoryA(Y)
-0.7  Features(Y,$F2) => CategoryB(Y)
1.1   Features(X,$F1) ^ Features(Y,$F2) => Link(X,Y)
+Inf  Link(X,Y) => CategoryA(X) ^ CategoryB(Y)
```

Ground-specific weights are computed by neural networks.
Infinite weights correspond to **hard** constraints.

Ground-Specific MLNs [Lippi & Frasconi, 2009]

4 Models and Frameworks

- Inference algorithms **do not change**
- Learning algorithms implement **gradient descent**

$$\frac{\partial P(y|x)}{\partial \theta_k} = \frac{\partial P(y|x)}{\partial w_i} \frac{\partial w_i}{\partial \theta_k}$$

where the **first** term is computed by MLN inference and the **second** term is computed by backpropagation

ProbLog [De Raedt et al., 2007]

4 Models and Frameworks

ProbLog is a **probabilistic extension of Prolog** where probabilities can be attached to ground facts or rules.

DeepProbLog extends ProbLog by computing such probabilities with neural networks in a framework for probabilistic reasoning

- Necessary to know $\text{Pro}(b)\text{Log}$
- Cannot (yet) perform collective classification

ProbLog [De Raedt et al., 2007]

4 Models and Frameworks

A **ProbLog** example (probabilistic logic program)

```
person(alice).  
person(bob).  
person(carl).  
movie(bladerunner).  
movie(titanic).  
friends(alice,bob).  
friends(bob,alice).  
likes(bob,bladerunner).  
  
0.2::friends(X,Y) :- person(X), person(Y).  
0.5::likes(X,M) :- person(X), movie(M).  
0.8::likes(Y,M) :- person(X), movie(M), likes(X,M), friends(X,Y).  
  
query(likes(_,_)).
```

Demo: <https://dtai.cs.kuleuven.be/problog/editor.html>

MLN vs. ProbLog

4 Models and Frameworks

Weights vs. probabilities

- In an MLN, the weight of formula F represents the log-odds between a world where F is true and a world where F is false, other things being equal
- In ProbLog we model directly the probability that a rule is true

DeepProbLog [Manhaeve et al., 2018]

4 Models and Frameworks

Query

 +  = ?

DeepProbLog Program

%Neural predicate

nn(net,[X],Y,[0..9]) :: digit(X,Y).

%Background knowledge

addition(X,Y,Z):- digit(X,N1),digit(Y,N2),
Z is N1+N2.

Logical Reasoning



addition(,,8) :- digit(,0),digit(,8),8 is 0+8.

...

addition(,,8) :- digit(,5),digit(,3),Z is 5+3.

...

Neural network evaluation

nn(net,[,Y],[0..9]) :: digit(,Y).

