

Fundamentals of Machine Learning:

Linear Regression, from Geometry to Maximum Likelihood

Prof. Andrew D. Bagdanov (andrew.bagdanov AT unifi.it)



UNIVERSITÀ
DEGLI STUDI
FIRENZE

DINFO
DIPARTIMENTO DI
INGEGNERIA
DELL'INFORMAZIONE

Introduction

Linear Regression as Geometric Curve Fitting

Linear Regression and Maximum Likelihood Estimation

Concluding remarks

Introduction

Lecture objectives

At the end of this lecture you will:

- Understand the **geometric** approach to linear regression problems as an application of **curve fitting**.
- Understand the **probabilistic** formulation of linear regression and know how the **Maximum Likelihood Estimate** of model parameters is derived.
- Recognize when **underfitting** and **overfitting** of models takes place.
- Understand how **regularization** can be used to **smoothly** limit the capacity of our models and mitigate **overfitting**.

Motivations

- Keep in mind and we work through the lecture today that we are heading towards a probabilistic model.
- We will start, however, with a straightforward geometric model that we will base out intuitions on.
- Interestingly, we will keep coming back to this "simple" model and verify that our intuitions are solid even from a probabilistic perspective.

Linear Regression as Geometric Curve Fitting

Linear Regression from the ground up

- We begin our study of **linear models for regression** from a ground-up, **geometric** perspective.
- We will define the **learning** problem as one of finding the optimal **parameters** \mathbf{w} of a **parametrized** hypothesis space \mathcal{H} .
- **Optimal** will be determined (for now) in terms of a simple, **geometric** error on a finite sample of **observations**.

A parametric, linear hypothesis space

- Consider first the class of univariate regression problems where we observe pairs of observations (x_i, t_i) , where:
 - $x_i \in \mathbb{R}$ are observations of the **independent** variable.
 - $t_i \in \mathbb{R}$ are observations of the **dependent** variable.
- We **assume** the dependent variable are related to the independent variable via a function f that **unknown** but **linear** in its parameters \mathbf{w} :

$$\begin{aligned}y(x \mid \mathbf{w}) &= w_0 + w_1 x \\ &= \mathbf{w}^T \begin{bmatrix} 1 \\ x \end{bmatrix}\end{aligned}$$

Quantifying the **goodness** of a solution

- OK, we have some **data** (i.e. observations of the independent and dependent variable).
- And we have a **hypothesis space** conveniently parameterized by \mathbf{w} .
- To actually **learn** something (i.e. **fit** the model to the data) we next need some sort of **loss** \mathcal{L} .
- This loss should measure the **badness** of an arbitrary hypothesis \mathbf{w}' – that is, **high** loss indicates a **bad** fit and **low** loss a **good** one.
- It will be a function of the **candidate** parameters and the **observed** data $\mathcal{D} = \{ (x_i, y_i) \mid i = 1, \dots, N \}$:

$$E(\mathbf{w} \mid \mathcal{D}) = \frac{1}{2} \sum_{i=1}^N [y(x_i \mid \mathbf{w}) - t_i]^2$$

Linear Regression: a Critique of Pure Reason

- We will look in more detail at this choice of this **loss**.
- But, for now let's consider the properties, good and bad, of this formulation.

$$E(\mathbf{w} \mid \mathcal{D}) = \sum_{i=1}^N [(w_0 + w_1 x_i) - t_i]^2$$

Linear Regression and Maximum Likelihood Estimation

Maximum Likelihood Estimation (MLE) and Least Squares

- Our **geometric** interpretation is an example of a **point estimate** of model parameters: it gives us a solution \mathbf{w}^* which is a **single point** in **parameter space**.
- We would like to move towards a **probabilistic** model of linear regression.
- This should, ideally:
 - Allow us to **reason** probabilistically about the quality of our solution.
 - Allow us to **quantify** belief in the quality of **individual** predictions.
 - Allow us to "bake in" **prior** knowledge about likely solutions.
- Our first step is the **Maximum Likelihood Estimate (MLE)** of the optimal parameters \mathbf{w} given the **observed data** \mathcal{D} .
- First, let's **beef up** our model a bit to allow **polynomial** functions of the input.

Linear basis function models

- The simplest **linear** model for regression just uses linear combinations of the input variables $\mathbf{x} = (x_1, \dots, x_D)^T$:

$$y(\mathbf{x} \mid \mathbf{w}) = w_0 + w_1x_1 + \dots + w_Dx_D$$

- This is a **linear** function in both \mathbf{w} (good) **and** \mathbf{x} (very limiting!).
- So, we can allow **linear** combinations **fixed** nonlinear functions of the input:

$$y(\mathbf{x} \mid \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

Linear basis function models

- w_0 is known as a **bias** parameter and allows for any fixed-offset in the output.
- It is often convenient to define a **dummy** basis function $\phi_0(\mathbf{x}) = 1$ so we can **compactly** write:

$$\begin{aligned}y(\mathbf{x} \mid \mathbf{w}) &= \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) \\ &= \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})\end{aligned}$$

- Where $\mathbf{w} = (w_0, \dots, w_{M-1})^T$ and $\boldsymbol{\phi} = (\phi_0, \dots, \phi_{M-1})^T$
- Common basis functions:

$\phi_i(x) = x^i$	$\phi_i = \exp\left\{-\frac{(x-\mu_j)^2}{2s^2}\right\}$	$\phi_i(x) = \tanh(x)$
Polynomial	Gaussian	Sigmoid

Maximum likelihood and least squares

- Neat, but where are our **probabilities** in all of this?!
- Let's go back to our original **assumption** that our **target** variable is the realization of a **deterministic** function $y(\mathbf{x}, \mathbf{w})$ with additive **Gaussian** noise:

$$t = y(\mathbf{x} \mid \mathbf{w}) + \varepsilon$$


- Where ε is a zero-mean, Gaussian random variable with **precision** β .
- Thus, we can write:

$$p(t \mid \mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t \mid y(\mathbf{x} \mid \mathbf{w}), \beta^{-1})$$

- Great! A **probability** distribution! Now let's learn something from data...

Maximum likelihood and least squares

- If you strip away all of the fancy mumbo jumbo, what we are doing is pretty literal and intuitive.
- At any point x_0 , our predictor qualifies its prediction by placing a Gaussian around it.



Maximum likelihood and least squares

- Now consider a dataset of **inputs** $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- Along with corresponding **target values** $\mathbf{t} = (t_1, \dots, t_N)^T$.
- Assuming that these samples are all identically and independently drawn from $p(t \mid \mathbf{x}, \mathbf{w}, \beta)$, we can write:

$$\begin{aligned} p(\mathbf{t} \mid \mathbf{X}, \mathbf{w}, \beta) &= \prod_{n=1}^N \mathcal{N}(t_n \mid y(\mathbf{x}_n, \mathbf{w}), \beta^{-1}) \\ &= \prod_{n=1}^N \mathcal{N}(t_n \mid \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \end{aligned}$$

- And we have a (very inconvenient) **likelihood** expression for our **data** under a specific model.

Maximum likelihood and least squares

- For **most** supervised learning problem we are not interested in modeling the distribution of **inputs**.
- Thus **X** will always appear in the conditioning variables and we will **omit** it for now to unclutter notation.
- Taking **logarithms** helps simplify the likelihood (thanks in part to the **exponential form** of the Gaussian):

$$\begin{aligned}\ln p(\mathbf{t} \mid \mathbf{w}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(t_n \mid \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} (\ln \beta - \ln(2\pi)) - \frac{\beta}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2\end{aligned}$$

Maximum likelihood and least squares

- Let's **maximize** this likelihood in \mathbf{w} . First the **gradient**:

$$\nabla \ln p(\mathbf{t} | \mathbf{w}, \beta) = \beta \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\} \phi(\mathbf{x}_n)^T$$

- And set it to **zero**:

$$0 = \sum_{n=1}^N t_n \phi(\mathbf{x}_n)^T - \mathbf{w}^T \left(\sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^T \right)$$

- Which after solving for \mathbf{w} gives us:

$$\mathbf{w}_{\text{ml}} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

Maximum likelihood and least squares


- This solution $(\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$ uses the **design matrix**:

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}$$

- An easy way to think of Φ :
 - A **row** is evaluation of **all basis functions** on the **corresponding training sample** (dimensionality M).
 - A **column** is the **corresponding basis function** evaluated on **all training samples** (dimensionality N).
- $\Phi^\dagger \equiv (\Phi^T \Phi)^{-1} \Phi^T$ is called the **Moore-Penrose Psuedo-Inverse**.

Regularized least squares


- So, what if we have data **distributed** like this:



- What might a good **basis** be? Can we recoup our **curve-fitting** approach?
- **Short answer:** Yes! We just saw an example of a polynomial basis!


Regularized least squares

- OK, let's go to town:




Regularized least squares

- OK, let's go to (quadratic) town:




Regularized least squares

- OK, let's go to (cubic) town:



Regularized least squares

- OK, let's **really** go to (20-degree) town:





Regularized least squares: over- and under-fitting

- Which is the **best** result? What does **best** mean here?
- We see in these plots examples of both **overfitting** (when M is too large) and **underfitting** (when M is too small).
- These results might seem counter intuitive: after all, a **degree-10** polynomial hypothesis space also **contains** all polynomials of **smaller** degrees.
- Why can't doesn't least squares find the **right** solution if we set $M = 100$ (for example)?
- The answer has to do with **model capacity** and the fact that the **noise** in our observations is easily captured by a **powerful** model.

Regularized least squares

- We can gain some insight by inspecting the **magnitude** of the parameters estimated by maximizing the likelihood.



	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Regularized least squares

- We will consider **regression objectives** of this form:

$$E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

- Where E_D will continue to be our least squares error with basis functions ϕ .
- The term E_W is called a **weight regularizer** (or just regularizer), and a very common one is the squared norm of the model weights:

$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

- So, the total error function to minimize becomes:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{w})\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

Regularized least squares

- This **regularized** error function:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{w})\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Is **still** a quadratic function in \mathbf{w} , and the optimal \mathbf{w} is:

$$\mathbf{w}_{\text{ML}} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- The Maximum Likelihood estimate of the β imprecision parameter can also be derived by maximizing $p(t \mid \mathbf{w}, \beta)$ wrt β :

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{i=1}^N [t_n - \mathbf{w}_{\text{ML}}^T \phi(\mathbf{x}_i)]^2$$


Live Demo!

- Let's take a look at how regularized least squares works in practice.
- [MAKE DEMO GO NOW]


Regularized least squares

- A more **general** regularized error is:


$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$




$q = 0.5$



$q = 1$



$q = 2$




$q = 4$

Regularized least squares

- A more **general** regularized error is:

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$



Concluding remarks

Good Old Linear Regression (tm)

- Today we took a fairly deep look at linear regression.
- Remember that linear is somewhat of a misnomer: anything that is linear in the model parameters, is a linear regression.
- In fact, we saw how flexible linear regression with basis functions is.
- So we can fit nonlinear models to data using pure linear regression.
- The Maximum Likelihood solution also has an analytical form, which is just extra nice.
- With this curve fitting model (splines are actually an instance of this) we can fit a lot of phenomena.

The way forward

- The purely geometric and MLE estimators we developed are examples of **point estimates**.
- They result in a **single point** in parameter space that is an **optimal** estimate (for some appropriate definition of "optimal").
- But, we cannot bring all of our **statistical** tools to bear and **quantify various types of belief**.
- In the next lecture we will develop a **Bayesian** theory of linear regression which addresses this issue.

Reading and Homework Assignments

Reading Assignment:

- Bishop: Chapter 3 (3.1, 3.2, 3.3)

Homework:

1. How does the **scale** of our input features affect the linear regression formulations we saw in this lecture? That is, if the scale of our input features and target is $\mathcal{O}(1)$ or $\mathcal{O}(1e6)$, does this make any difference?
2. Does the scale of our input features matter more (or less) if we use **regularized** least squares?
3. Why would it be problematic if, when using a K -degree polynomial basis for regression, the design matrix Φ is **rank deficient** (i.e. $\text{rank}(\Phi) < K$)?