

Fundamentals of Machine Learning:

Nonparametric Methods: Nearest Neighbors and Kernel Density Estimation

Prof. Andrew D. Bagdanov ([andrew.bagdanov AT unifi.it](mailto:andrew.bagdanov@unifi.it))



UNIVERSITÀ
DEGLI STUDI
FIRENZE

DINFO
DIPARTIMENTO DI
INGEGNERIA
DELL'INFORMAZIONE

Outline

Introduction

Leftovers: Classifier Evaluation

Density estimation: Histograms Revisited

Density Estimation: Kernel Density Estimators

Nadaraya-Watson Kernel Regression

K-Nearest Neighbors

Concluding Remarks

Introduction

Motivations

- We have seen a **variety** of methods for classification and regression.
- So far, all of them are what is known as **parametric** methods.
- These methods concentrate on **learning** by estimating the optimal model **parameters** (in some way).
- The resulting model is typically **simple** (at least so far):

$$f(\mathbf{x}; \mathbf{w}, b) = \mathbf{w}^T \mathbf{x} + b$$

- A characteristic of these techniques is that the "**knowledge**" learned from the data is completely encapsulated in the **parameters**.
- This is sometimes called a **bottleneck** and can help generalization.

Lecture objectives

After this lecture you will:

- Understand how to **evaluate** classifier performance.
- Understand the basics of **kernel density estimators** and how to use them to model probability density functions.
- Understand how **kernel density estimation** approach can be applied to regression problems (e.g. **Nadaraya-Watson**).
- Understand how the **Nearest-neighbor Algorithm** works in theory and in practice.

Leftovers: Classifier Evaluation

Evaluating Classifiers: Balanced Problems



- How do we measure the **performance** of trained classifiers?
- If you have a **balanced** classification problem (like most we will see in the labs), use the classifier **accuracy**:

$$\text{accuracy} = \frac{\# \text{ correctly classified test samples}}{\# \text{ test samples}}$$

- Why might this be a **bad metric** if the problem is unbalanced?

Evaluating Classifiers: Confusion Matrices

- A good tool to get an overview of the **types** of errors a classifier is making is the **confusion matrix**:



Evaluating Classifiers: Unbalanced Problems

- We begin by dissecting the classifications:
 - True positives (TP): we predicted yes, and sample does belong to class.
 - True negatives (TN): we predicted no, and the sample does not belong to class.
 - False positives (FP): we predicted yes, but the sample does not belong to class – also known as a Type I error.
 - False negatives (FN): we predicted no, but sample does belong to the class – also known as a Type II error.
- We can then define some useful metrics for unbalanced problems:


$$\text{Precision}(c) = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall}(c) = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{F1}(c) = 2 \frac{\text{Precision}(c) * \text{Recall}(c)}{\text{Precision}(c) + \text{Recall}(c)}$$

Evaluating Classifiers: Unbalanced Problems

- We usually understand precision and recall using Precision-recall (PR) and/or Receiver-Operating Characteristic (ROC) curves:



Evaluating Classifiers: Analysis

- The **evaluation metric** to use is usually clear given the **type** of problem (classification, regression, retrieval) and the **distribution** of training data (balanced, unbalanced).
- Getting a **reliable** estimate of a classifier can be tricky as it clearly depends on your train/test split.
- This becomes the central issue when performing **model selection** via cross-validation..
- In **sklearn: sklearn.metrics**

Density estimation: Histograms Revisited

The problem with parametric models

- The **probabilistic** models we have studied thus far are based on **parametric** models.
- We estimate a **small** number of distribution parameters (e.g. mean and covariance of **Gaussians**).
- Such a **prior** assumption of a distributional form is a limitation of these approaches.
- **Especially** when these assumptions turn out to be **inaccurate**.

Idea: can we **dispense** with the parametric assumption, and just let the **data** speak for itself?

Histograms: piecewise constant density estimators


- A **histogram** partitions the domain into **fixed-width** bins.
- We then **count** the number of observations that fall into each discrete bin.
- We can turn this representation into a **probability density** by dividing by the total number of observations N and the bin width Δ_i :

$$p_i = \frac{n_i}{N\Delta_i}$$

- This yields a density $p(x)$ that is **constant** over the width of each bin (usually $\Delta_i = \Delta$).


Histograms: piecewise constant density estimators

- This gives a **density estimator** with a single hyperparameter:



Histograms: piecewise constant density estimators

- Selection of Δ is critical – too small and the estimator has **high variance**.
- Too large and the curve is **too smooth**.
- The model is **discontinuous** at bin boundaries.
- The data are **not** needed at test time, but histograms scale very poorly with dimension (M^D).



Histograms: piecewise constant density estimators

- Despite their limitations, histograms tell us something important about **density estimation**.
- To estimate the density at some point x in the domain, we should look at data points that lie **near** x .
- This use of **locality** assumes we have some meaningful distance metric in the domain (e.g. **Euclidean**).
- Let's look at some common **nonparametric methods** used in Machine Learning.

Density Estimation: Kernel Density Estimators


A binomial aside

- To motivate how we can extend our **histogram** idea to a purely **local** estimator, we need to consider the **binomial distribution**.
- The binomial distribution represents the probability of seeing m **successes** in a sequence of N **Bernoulli trials**:

$$\begin{aligned}\text{Bin}(m|N, \mu) &= \binom{N}{m} \mu^m (1 - \mu)^{(N-m)} \\ \mathbb{E}(m) &= \sum_{m=0}^{N} m \text{Bin}(m|N, \mu) \\ &= N\mu\end{aligned}$$

A binomial aside

- The binomial is **unimodal** and somewhat **Gaussian-like**:



Generalizing

- Assume we have observations drawn from some **unknown** density $p(\mathbf{x})$ in a D -dimensional Euclidean space.
- Similar to a **histogram**, we consider a small region \mathcal{R} around \mathbf{x} .
- The probability mass associated with \mathcal{R} is:

$$P = \int_{\mathcal{R}} p(\mathbf{x}) d\mathbf{x}$$

- If we have N observations drawn from $p(\mathbf{x})$, each one has probability P of falling in \mathcal{R} .
- So, the **total number of points** K falling in \mathcal{R} will be **binomially** distributed:

$$\text{Bin}(K|N, P) = \binom{N}{K} P^K (1 - P)^{(N-K)}$$

Generalizing

- Using the formulas for **mean** and **variance** of the Binomial we can analyze the **fraction** of points falling in \mathcal{R} :

$$\begin{aligned}\mathbb{E}(K/N) &= P \\ \text{var}(K/N) &= \frac{P(1 - P)}{N}\end{aligned}$$

- For large N this distribution will be **concentrated** around the P and so:

$$K \approx NP$$

Generalizing

- For large N this distribution will be **concentrated** around the P and so:

$$K \approx NP$$

- On the other hand, we can assume \mathcal{R} with **volume** V is sufficiently small (and hence $p(\mathbf{x})$ is constant) and so:

$$P \approx p(\mathbf{x})V$$

- Leading us to:

$$p(\mathbf{x}) \approx \frac{K}{NV}$$

Generalizing

- Note that this estimate of $p(x)$ is based on two contradictory assumptions:
 - That \mathcal{R} is sufficiently **small** and thus $p(x)$ is **constant** in it.
 - That \mathcal{R} is sufficiently **large** so that the number of points falling into it are enough to **peak** the binomial.
- This leaves us with **two** choices:
 1. Fix K and let the **data** determine what V should be (**Nearest Neighbors**).
 2. Fix V and let the **data** determine what K should be (**Kernel Density Estimation**).

Kernel Density Estimation

- Let's take \mathcal{R} to be a **unit hypercube** centered at the origin.
- This is conveniently described by the **kernel function**:

$$k(\mathbf{u}) = \begin{cases} 1 & \text{if } |u_i| \leq 1/2, \text{ for } i = 1, \dots, D \\ 0 & \text{otherwise} \end{cases}$$

- So, for any point \mathbf{x} , $k((\mathbf{x} - \mathbf{x}_n)/h) = 1$ if the point \mathbf{x}_n lies within a **hypercube** with size h centered on \mathbf{x} .
- The total number of points K lying in such a hypercube around \mathbf{x} is then:

$$K = \sum_{n=1}^N k\left(\frac{\mathbf{x} - \mathbf{x}_n}{h}\right)$$

Kernel Density Estimation

- Going back to our equation for $p(\mathbf{x})$, we have:

$$\begin{aligned} p(\mathbf{x}) &= \frac{K}{NV} \\ &= \frac{1}{N} \sum_{n=1}^N \frac{1}{h^D} k\left(\frac{\mathbf{x} - \mathbf{x}_n}{h}\right) \end{aligned}$$

- This still suffers from discontinuities at the **hypercube boundaries**, so a more common kernel is the **Gaussian**:

$$p(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N \frac{1}{(2\pi h^2)^{D/2}} \exp\left\{-\frac{||\mathbf{x} - \mathbf{x}_n||^2}{2h^2}\right\}$$


Kernel Density Estimation

- This is called a **Kernel (or Parzen) Density Estimator**.
- The hyperparameter h is called the **kernel bandwidth** that controls the amount of **smoothing** that is performed.
- Kernels functions should satisfy:

$$k(u) \geq 0, \int uk(u)du = 0, \text{ and } \int k(u)du = 1$$

- These type of approaches are often called **local methods** because to estimate $p(x)$ they look at **data close** to x .

Kernel Density Estimation



Nadaraya-Watson Kernel Regression

Locality in x and y

- We can apply this type of local method to regression problems.
- Recall that an optimal predictor for regression is:

$$\begin{aligned}\mathbb{E}(t | \mathbf{x}) &= \int tp(t | \mathbf{x})dt \\ &= \int t \frac{p(\mathbf{x}, t)}{p(\mathbf{x})} dt\end{aligned}$$

- Now we use Kernel Density Estimates as approximations:


$$\hat{p}(\mathbf{x}, t) = \frac{1}{N} \sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n) k_h(t - t_n) \quad \hat{p}(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n)$$

Nadaraya-Watson: locally weighted regression

- Our approximation of the optimal estimator is now:

$$\begin{aligned}\mathbb{E}(t | \mathbf{x}) &= \int t \frac{p(\mathbf{x}, t)}{p(\mathbf{x})} dt \\ &= \int t \frac{\frac{1}{N} \sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n) k_h(t - t_n)}{\frac{1}{N} \sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n)} dt \\ &= \frac{\sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n) \cancel{\int t k_h(t - t_n) dt}}{\sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n)} \\ &= \frac{\sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n) \textcolor{red}{t_n}}{\sum_{n=1}^N k_h(\mathbf{x} - \mathbf{x}_n)}\end{aligned}$$

An approximately optimal estimator



K-Nearest Neighbors

The problem of h


- A problem with Kernel Density Estimation is that the bandwidth h is fixed for all kernels.
- If we have parts of the space with a lot of samples, this can lead to oversmoothing and loss of detail if h is too large.
- If we reduce h , in low-density regions we will have noisy estimates.
- What if, instead of fixing V (via the bandwidth) and deriving an expression for K , we fix K and go in the other direction.

K-Nearest Neighbors

- The K-Nearest Neighbors approach is algorithmic:
 - We fix K and consider a hypersphere around a point \mathbf{x} where we want to estimate $p(\mathbf{x})$.
 - We then increase the radius of this hypersphere until exactly K points from the training set are inside it.
 - We then compute $p(\mathbf{x}) = \frac{K}{NV}$ for V equal to the volume of the resulting hypersphere.

K-Nearest Neighbors

- This approach is **discontinuous** and **noisy** for density estimation:



K-Nearest Neighbors Classification

- The K-Nearest Neighbors (KNN) approach is a **very** convenient method for **classification**, however.
- Intuitively, we make a KNN density estimate for each **class conditional** density, and then apply **Bayes' Rule**.
- Suppose we have N_k examples from each class k (so $\sum_k N_k = N$).
- To classify a new point \mathbf{x} we do the **hypersphere trick** assuming the volume needed around \mathbf{x} to include K points is $V_{\mathbf{x}}$.
- If in this hypersphere we have K_k points from class k :

$$p(\mathbf{x} | \mathcal{C}_k) = \frac{K_k}{N_k V_{\mathbf{x}}}$$



- The **marginal density** and class priors are similarly estimated:

$$p(\mathbf{x}) = \frac{K}{NV_{\mathbf{x}}} \quad p(\mathcal{C}_k) = \frac{N_k}{N}$$

K-Nearest Neighbors Classification

- Enter Bayes:

$$p(\mathcal{C}_k \mid \mathbf{x}) = \frac{p(\mathbf{x} \mid \mathcal{C}_k)p(\mathcal{C}_k)}{p(\mathbf{x})} = \frac{K_k}{K}$$




Concluding Remarks

Local methods

- In this lecture we have seen several **nonparametric** and **local** methods for:
 1. Estimating densities (**Kernel Regression**).
 2. Regression (**Nadaraya-Watson**)
 3. Classification (**K-Nearest Neighbors**).
- These techniques are **simple** and **intuitive** and **explainable**.
- They have a small (but **important!**) number of hyperparameters: the **kernel bandwidth** h (often more than one), or the number of **neighbors** K .
- These methods have **nice** convergence properties (e.g. K-Nearest Neighbors is **near-optimal** for $K = 1$ as $N \rightarrow \infty$).
- These methods however have a **huge** practical drawback?
Has anyone seen it yet?

Space inefficiency

- To make predictions, all these methods *require that all training data be available at inference time.*
- In practice, we typically use a **data structure** (e.g. a **KD-tree** or **Ball-Tree**) to **efficiently** and **approximately** retrieve the K points in \mathcal{D} closest to a point x .



Reading and Homework Assignments

Reading Assignment:

- Bishop: Chapter 2 (2.5), Chapter 6 (6.3)
- Scikit-learn User Guide (sections 2.8, 1.3, 1.6)

Homework:

1. Try a K-Nearest Neighbor classifier on the classification problems (synthetic and real) from the lab on Classification. Visualize the decision boundary and observe the qualitative differences (with varying K between it and that of an SVM or a Bayes generative classifier).
2. Implement the Nadaraya-Watson Kernel Regressor with a Gaussian Kernel. Use it to estimate a regressor for the example in the Linear Regression laboratory. How should the bandwidth parameter h be tuned?