

Fundamentals of Machine Learning:

Unsupervised Learning I: K-Means Clustering and Gaussian Mixture Models

Prof. Andrew D. Bagdanov ([andrew.bagdanov AT unifi.it](mailto:andrew.bagdanov@unifi.it))



UNIVERSITÀ
DEGLI STUDI
FIRENZE

DINFO
DIPARTIMENTO DI
INGEGNERIA
DELL'INFORMAZIONE

Outline

Introduction

K-means Clustering

SKIPPABLE Gaussian Mixture Models

Concluding Remarks

Introduction

Motivations

- In the previous lecture we saw how **kernel density estimators** (or **Parzen** density estimation) can be used to infer a marginal data distribution.
- This is already an example of an **unsupervised** learning problem.
- We are given **data** without any **targets** and we should learn something "useful" from this data.
- In this lecture we will see two more standard approaches to **unsupervised learning**:
 1. The **K-Means** algorithm, which learns K **clusters** of data in Euclidean spaces.
 2. The **Gaussian Mixture Model**, which fits a generative parametric **mixture** of K Gaussians to the data.

Lecture objectives

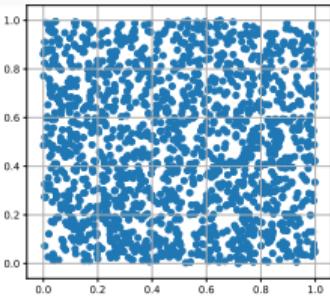
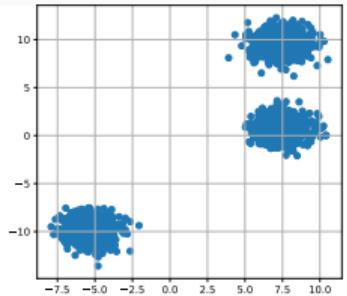
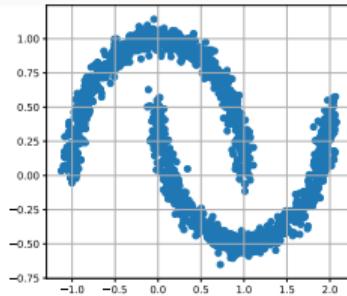
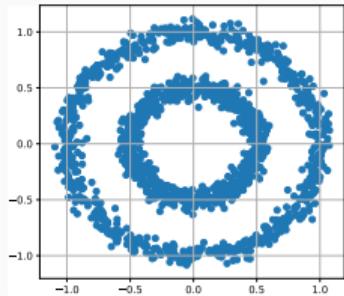
After this lecture you will:

- Understand how the K-Means algorithm **partitions** the input data by assigning each point to **exactly** one cluster.
- Understand how Gaussian Mixture Models represent the **data likelihood** as **superposition** of Gaussian distributions.
- Understand how **latent variables** can be **entangled** with observed variables to render tractable certain optimization problems.
- Understand how **Expectation Maximization (EM)** is used to learn the parameters of a **GMM** from data.

K-means Clustering

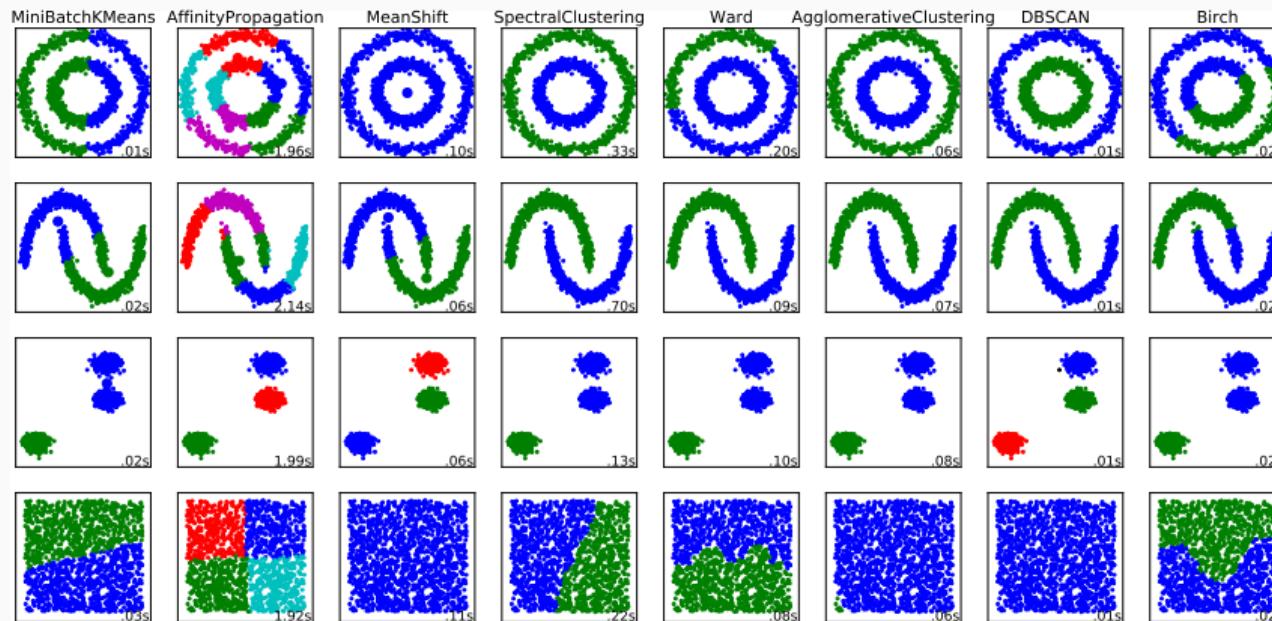
Clustering: Motivation

- What if you have data with **structure**, but you don't know what this structure is or have any a priori model for it?



Clustering: Motivation (continued)

- Clustering refers to techniques that learn groups of related data points in feature space:



Clustering: The k-Means Algorithm

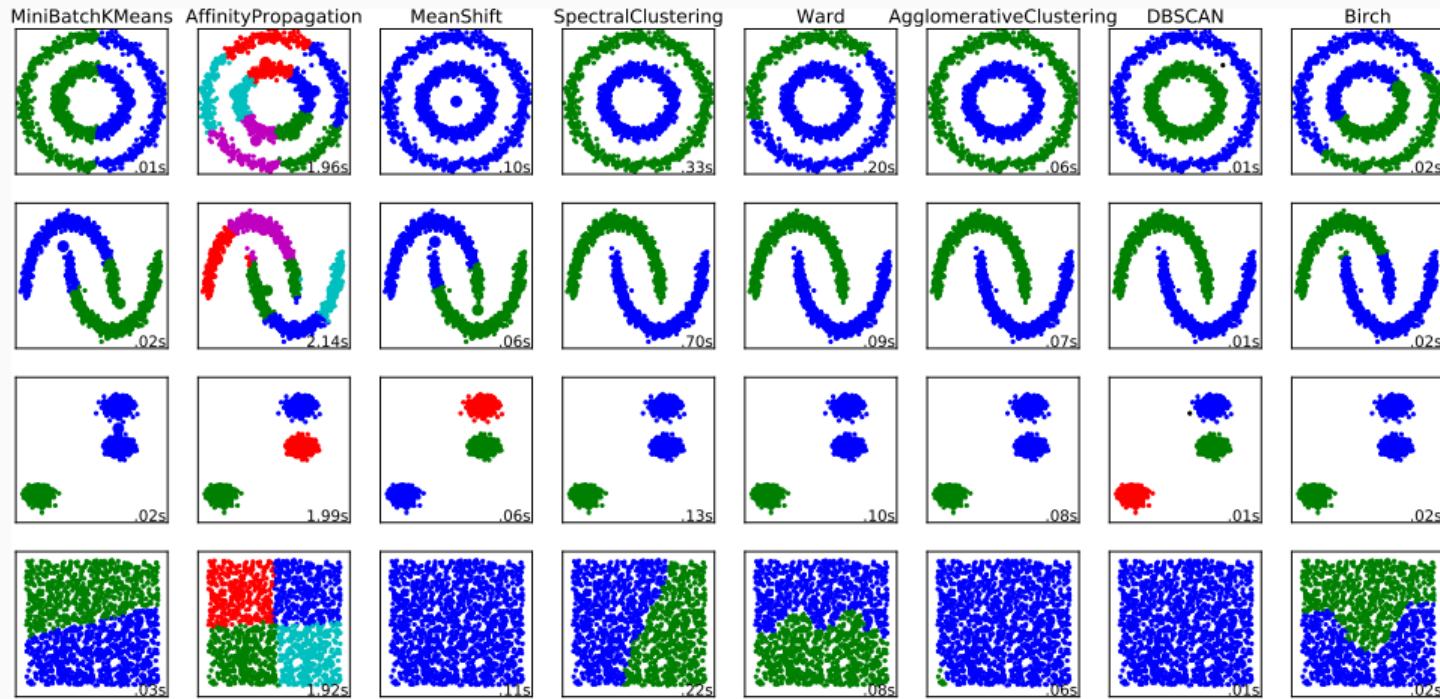
- **k-Means** is a very simple algorithm that associates each data point with one of k **means** or **centroids** in the data space:
- Given an initial set of k means $\mathbf{m}_1^1, \dots, \mathbf{m}_k^1$, the alternates between two steps:
 1. **Assignment**: $S_i^t = \{ \mathbf{x}_p \mid \|\mathbf{x}_p - \mathbf{m}_i^t\| \leq \|\mathbf{x}_p - \mathbf{m}_j^t\| \forall j \}$.
 2. **Update**: $\mathbf{m}_i^{t+1} = \frac{1}{|S_i^t|} \sum_{\mathbf{x}_i \in S_i^t} \mathbf{x}_i$.
- The algorithm **converges** when the cluster assignments no longer change.
- **Note**: this algorithm is **not** guaranteed to find the **global optimum clusters**.

Clustering: Analysis

- Clustering is a robust technique (in that it **never fails**).
- Its usefulness depends on the **data distribution** and which type of clustering you perform.
- In **sklearn**:
 - `cluster.AffinityPropagation`
 - `cluster.AgglomerativeClustering`
 - `cluster.Birch`
 - `cluster.DBSCAN`
 - `cluster.KMeans`
 - `cluster.MeanShift`
 - `cluster.SpectralClustering`
- See **sklearn.cluster** for more details.

Clustering: Motivation (continued)

- It's useful to **revisit** this plot:



The K-Means Algorithm

- Assume we have a dataset $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ of N samples of a D -dimensional Euclidean variable \mathbf{x} .
- Goal: partition \mathcal{D} into K clusters.
- What is a cluster? Intuitively, a subset of \mathcal{D} is a cluster if the sum of inter-point distances in it are small.
- To capture this intuition, we introduce a set of K D -dimensional prototype vectors $\boldsymbol{\mu}_k$ for $k = 1, \dots, K$.
- Then, for each sample \mathbf{x}_n we introduce a binary indicator variable r_{nk} :

$$r_{nk} = \begin{cases} 1 & \text{if } \mathbf{x}_n \text{ is assigned to cluster } k \\ 0 & \text{otherwise} \end{cases}$$

The K-Means Algorithm

- Now we can define an **objective function**:

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|\mathbf{x}_n - \boldsymbol{\mu}_k\|^2$$

- This results in an optimization problem in r_{nk} and $\boldsymbol{\mu}_k$.
- We can approach this with an **alternating** iterative algorithm which:
 - (E-Step): Minimizes J with respect to r_{nk} keeping $\boldsymbol{\mu}_k$ fixed.
 - (M-Step): Minimizes J with respect to $\boldsymbol{\mu}_k$ keeping r_{nk} fixed.

K-Means: The E-Step

- Consider holding the μ_k fixed and finding the optimal r_{nk} .
- J is a linear function of r_{nk} , and we can optimize for each n independently.
- So our rule is to just assign x_n to the closest μ_k :

$$r_{nk} = \begin{cases} 1 & \text{if } k = \arg \min_j \|x_n - \mu_j\|^2 \\ 0 & \text{otherwise} \end{cases}$$

K-Means: The M-Step

- Now consider holding the r_{nk} fixed and minimizing J in μ_k .
- J is quadratic in μ_k , so we set the gradient wrt them to zero.

$$2 \sum_{n=1}^N r_{nk} (\mathbf{x}_n - \boldsymbol{\mu}_k) = 0 \quad \forall k$$

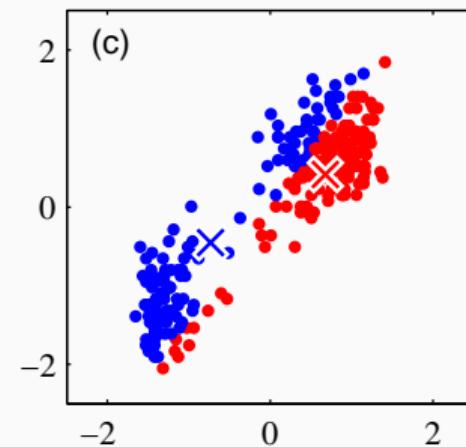
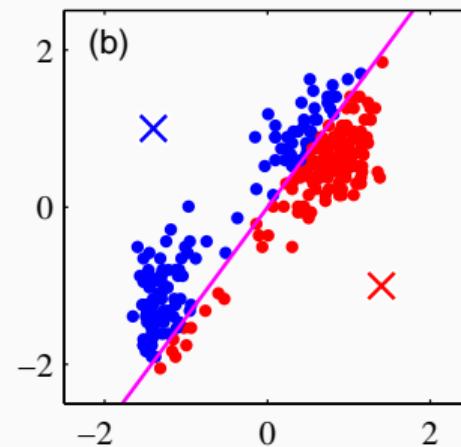
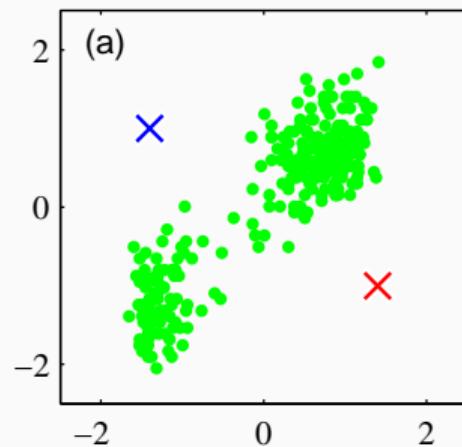
- Which is easy to solve for μ_k :

$$\boldsymbol{\mu}_k = \frac{\sum_{n=1}^N r_{nk} \mathbf{x}_n}{\sum_{n=1}^N r_{nk}}$$

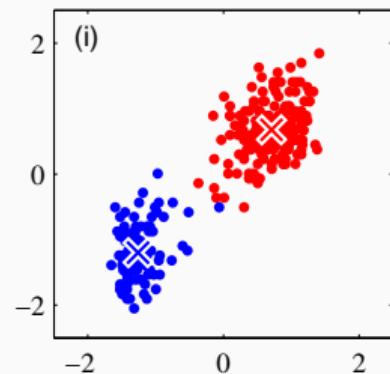
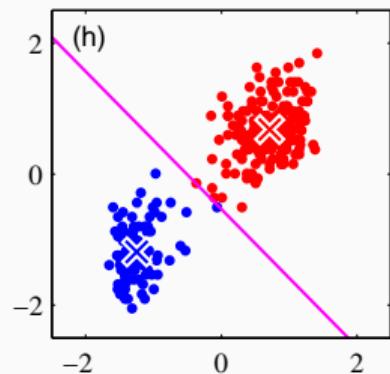
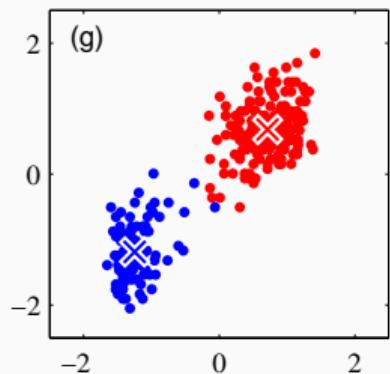
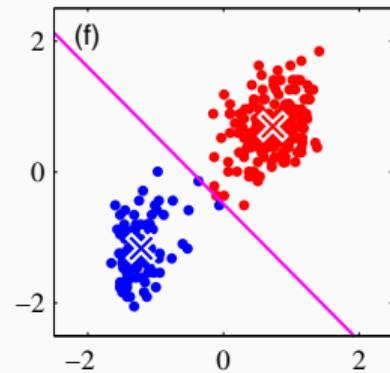
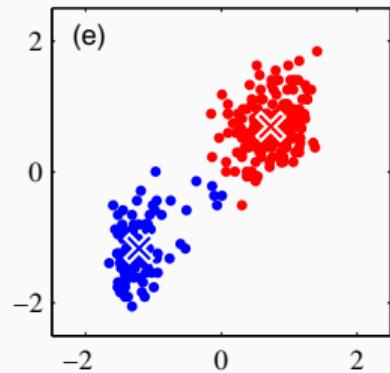
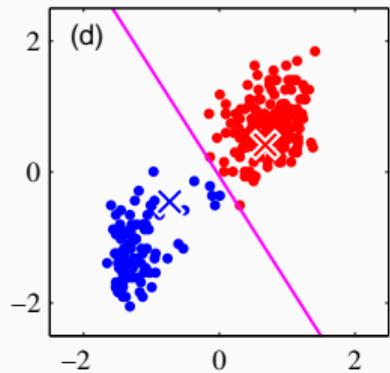
- And even easier to interpret.

K-Means: Analysis

- The two steps of the K-Means algorithm are alternated until the assignments do not change in the E-Step.
- Each step of K-Means reduces the value of J , convergence is guaranteed.
- Depending on initialization, however, K-Means can converge to a bad local minimum.



K-Means: Analysis

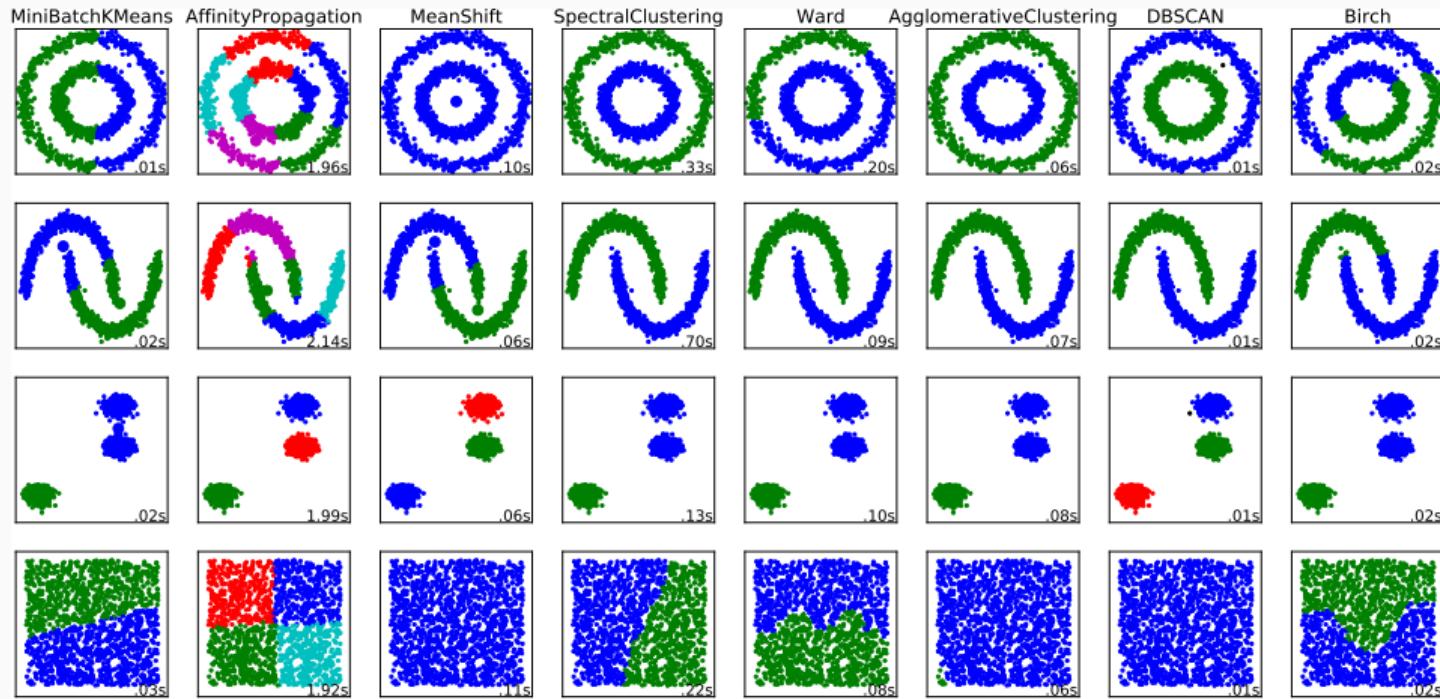


Clustering: Analysis

- Clustering is a robust technique (in that it **never fails**).
- Its usefulness depends on the **data distribution** and which type of clustering you perform.
- In **sklearn**:
 - `cluster.AffinityPropagation`
 - `cluster.AgglomerativeClustering`
 - `cluster.Birch`
 - `cluster.DBSCAN`
 - `cluster.KMeans`
 - `cluster.MeanShift`
 - `cluster.SpectralClustering`
- See **sklearn.cluster** for more details.

Clustering: Analysis

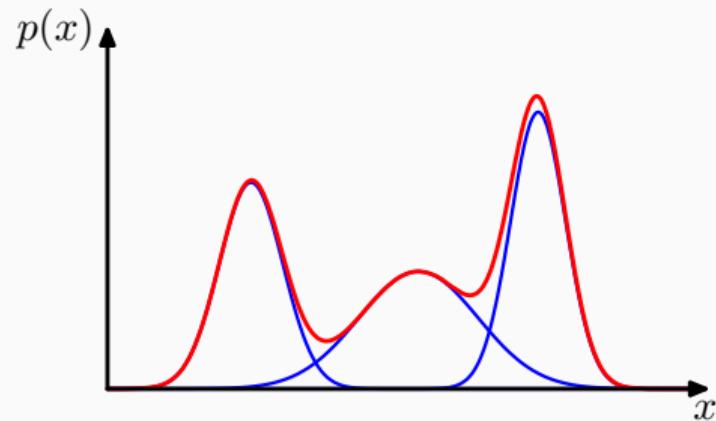
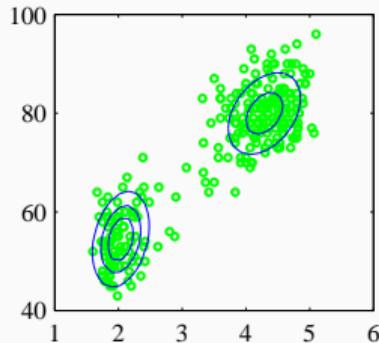
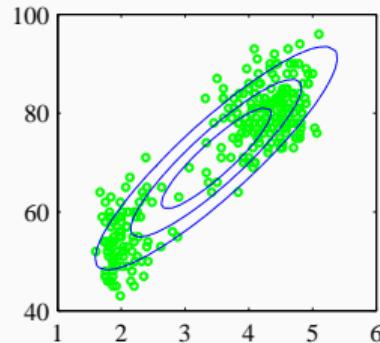
- It's useful to **revisit** this plot:



SKIPPABLE Gaussian Mixture Models

Gaussian Mixtures Models (GMMs) and Expectation Maximization (EM)

- Modeling distributions with a **single** Gaussian can work well in practice.
- However, sometimes it is a **terrible** approximation of reality.



GMMs: a Latent Variable Model

- We have already seen the **kernel density estimation** approach – slap a **Gaussian** around every point and **normalize**.
- In these examples, though, it seems clear that the data are generated from some **finite** and **fixed** number of generating distributions.
- We can instead fit a parametric **Mixture of Gaussians** – a **Gaussian Mixture Model (GMM)** to data:

$$p(\mathbf{x}; \boldsymbol{\theta}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- This is a **superposition** of K Gaussians, each of which is called a **component** of the mixture with **mixing coefficient** π_k .
- We require $0 \leq \pi_k \leq 1$ and $\sum_k \pi_k = 1$. **Why?**

GMMs: a Latent Variable Model

- OK, but if we try to use Maximum Likelihood to estimate μ_k and Σ_k :

$$\begin{aligned} p(\mathcal{D}; \theta) &= \prod_{n=1}^N p(\mathbf{x}_n | \theta) \\ &= \prod_{n=1}^N \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \end{aligned}$$

- Taking the logarithm (which won't change the **maximum**):

$$\ln p(\mathcal{D}; \theta) = \sum_{n=1}^N \ln \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}, \boldsymbol{\Sigma}_k)$$

- And we run into a dead end... **Why?**

GMMs: a Latent Variable Model

- Let's look at the mixture equation again:

$$p(\mathbf{x}; \boldsymbol{\theta}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- This type of **superposition** is difficult to fit to data without extra information about which **components** are “responsible” for specific points.
- So, we introduce a K -dimensional discrete **latent** random variable \mathbf{z} having a one-hot representation.
- A **latent** variable is one that is never actually **observed**, but that we introduce to **simplify** the optimization problem (eventually).

GMMs: a Latent Variable Model

- Define \mathbf{z} to be a one-hot random variable of dimension K :
 $z_k \in \{0, 1\}$ and $\sum_k z_k = 1$.
- We will now entangle this variable with \mathbf{x} by defining the joint distribution $p(\mathbf{x}, \mathbf{z})$ in terms of the marginal $p(\mathbf{z})$ and the conditional distributions $p(\mathbf{x} | \mathbf{z})$.
- The marginal distribution $p(\mathbf{z})$ is specified in terms of the mixing coefficients:

$$p(z_k = 1) = \pi_k \quad p(\mathbf{z}) = \prod_{k=1}^K \pi_k^{z_k}$$

- The conditional distribution of \mathbf{x} given a specific \mathbf{z} is similarly:

$$p(\mathbf{x} | \mathbf{z}) = \prod_{k=1}^K \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)^{z_k}$$

GMMs: a Latent Variable Model

- The joint distribution over \mathbf{x} and \mathbf{z} is $p(\mathbf{x}, \mathbf{z}) = p(\mathbf{z})p(\mathbf{x} | \mathbf{z})$.
- If we **marginalize** \mathbf{z} away, we get our mixture back:

$$\begin{aligned} p(\mathbf{x}) &= \sum_{\mathbf{z}} p(\mathbf{z})p(\mathbf{x} | \mathbf{z}) \\ &= \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \end{aligned}$$

- One might rightly ask: **so what?**
- The point is that this allows us to work with the **joint** distribution instead of the **unstructured** marginal.
- Because of this representation of $p(\mathbf{x})$, if we have N observations $\mathbf{x}_1, \dots, \mathbf{x}_N$, we will have a latent variable \mathbf{z}_n for each \mathbf{x}_n .

GMMs: a Latent Variable Model

- This is an important quantity in what follows:

$$\begin{aligned}\gamma(z_k) \equiv p(z_k = 1 | \mathbf{x}) &= \frac{p(z_k = 1)p(\mathbf{x} | z_k = 1)}{\sum_{k=1}^K p(z_k = 1)p(\mathbf{x} | z_k = 1)} \\ &= \frac{\pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}\end{aligned}$$

- π_k is the **prior** probability of $z_k = 1$, and $\gamma(z_k)$ is the corresponding **posterior**.
- We can think of $\gamma(z_k)$ as the **responsibility** component k takes for **explaining** observation \mathbf{x} .

GMMs: Fitting the Model

- Again, we have a **dataset** $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ that we want to fit a GMM to.
- We can represent \mathcal{D} as an $N \times D$ matrix \mathbf{X} .
- The corresponding **latent variables** will be an $N \times K$ matrix \mathbf{Z} .
- Assuming the \mathbf{x}_n are i.i.d. we can "just" optimize:

$$\ln p(\mathbf{X} | \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right\}$$

- We know how to do this: take **derivatives** wrt $\boldsymbol{\mu}$, $\boldsymbol{\Sigma}$, and $\boldsymbol{\pi}$ and set them to **0**.

GMMs: Fitting the Model

- Our **data likelihood**:

$$\ln p(\mathbf{X} \mid \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \right\}$$

- The **gradient** wrt $\boldsymbol{\mu}_k$:

$$0 = \sum_{n=1}^N \underbrace{\frac{\pi_k \mathcal{N}(\mathbf{x}_n \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_j \pi_j \mathcal{N}(\mathbf{x}_n \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}}_{=\gamma(z_{nk})} \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k)$$

- Defining $N_k = \sum_{n=1}^N \gamma(z_{nk})$ (**how to interpret this?**):

$$\boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) \mathbf{x}_n$$

GMMs: Fitting the Model

- By a similar (but longer) line of reasoning:

$$\boldsymbol{\Sigma}_k = \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk})(\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^T$$

- For the mixing coefficients we must constrain them to sum to 1:

$$\ln p(\mathbf{X} | \boldsymbol{\pi}, \boldsymbol{\mu}, \boldsymbol{\Sigma}) + \lambda \left(\sum_{k=1}^K \pi_k - 1 \right)$$

- The partials wrt π_k are:

$$\begin{aligned} 0 &= \sum_{n=1}^N \frac{\mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_j \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)} + \lambda \\ &= \frac{1}{\pi_k} \sum_{n=1}^N \frac{\pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_j \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)} + \lambda \end{aligned}$$

GMMs: Fitting the Model

- Again, our responsibilities appear (after multiplying by 1!):

$$\begin{aligned} 0 &= \frac{1}{\pi_k} \sum_{n=1}^N \frac{\pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_j \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)} + \lambda \\ &= \frac{N_k}{\pi_k} + \lambda \end{aligned}$$

- The partials wrt our Lagrange multiplier:

$$0 = \sum_{k=1}^K \pi_k - 1 \implies \sum_{k=1}^K \pi_k = 1 \quad (\text{duh})$$

- Combining, we have:

$$\pi_k = \frac{N_k}{N}$$

Expectation Maximization (EM) for GMMs

1. Initialize π , μ , and Σ .
2. **E step:** Compute responsibilities for current parameters.

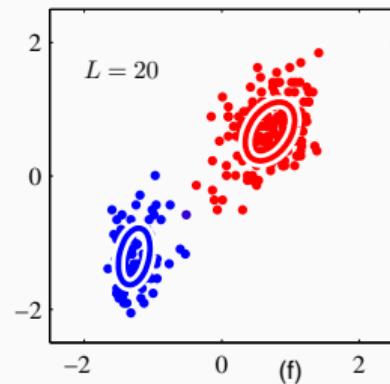
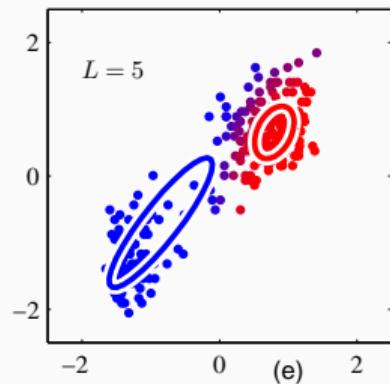
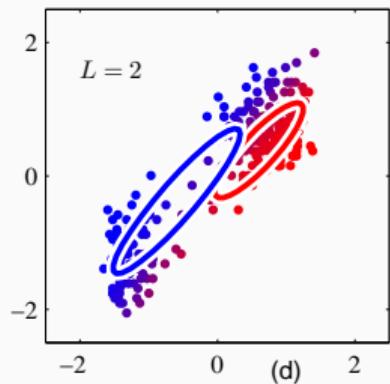
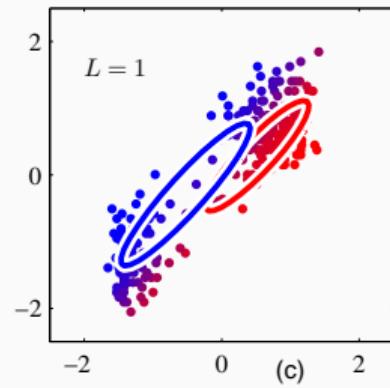
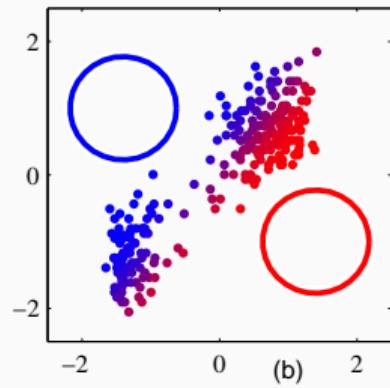
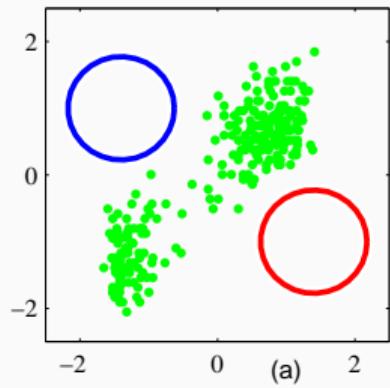
$$\gamma(z_{nk}) = \frac{\pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}, \quad N_k = \sum_{n=1}^N \gamma(z_{nk})$$

3. **M step:** Re-estimate parameters with current $\gamma(z_{nk})$ (to **maximize** likelihood).

$$\begin{aligned}\boldsymbol{\mu}_k^{\text{new}} &= \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) \mathbf{x}_n \\ \boldsymbol{\Sigma}_k^{\text{new}} &= \frac{1}{N_k} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu}_k^{\text{new}})(\mathbf{x}_n - \boldsymbol{\mu}_k^{\text{new}})^T \\ \pi_k^{\text{new}} &= \frac{N_k}{N}\end{aligned}$$

4. Check for **convergence** and (maybe) **repeat**.

Visualizing EM



Concluding Remarks

Clustering

- Clustering in general – and K-Means in particular – are **robust** methods for unsupervised learning.
- Clustering is used in a **broad variety of applications**.
- It is also useful to understand what, if any, **structure** exists in your data.
- There are many types of clustering algorithms, each of which has its own **affordances**.

Gaussian Mixtures and the EM Algorithm

- Gaussian Mixture Models can be thought of as a type of **soft** K-Means.
- In this lecture we have seen how introducing **latent** variables can (eventually) simplify optimization of otherwise **intractable** objectives.
- **GMMs** also have a natural **generative** interpretation:
 1. **Sample** a mixture component k according to $p(\pi)$.
 2. Then **sample** a point from $\mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$.
- Note that the **Expectation Maximization (EM)** algorithm is **broadly** applicable to estimation problems of joint distributions over **observed** and **latent** variables.
- Here we have only seen **one** such application of EM.

Reading and Homework Assignments

Reading Assignment:

- Bishop: Chapter 2 (2.3.9), Chapter 9 (9.1, 9.2)
- Scikit-learn User Guide (sections 2.1, 2.3)