

Fundamentals of Machine Learning:

Ensemble Models

Prof. Andrew D. Bagdanov ([andrew.bagdanov AT unifi.it](mailto:andrew.bagdanov@unifi.it))



UNIVERSITÀ
DEGLI STUDI
FIRENZE
DINFO
DIPARTIMENTO DI
INGEGNERIA
DELL'INFORMAZIONE

Outline

Introduction

Committees and Bagging

SKIPPABLE Boosting

Tree Models

Concluding Remarks

Introduction

The power of consensus

- All of the **regression** and **classification** models we have considered thus far are **singular** – we fit **one** model, and predict using **one** model.
- Often we can **improve** over singular models by **combining** multiple ones.
- We might train M models, and then take the **average** of their predictions.
- Or, we could train models **sequentially** and encourage **subsequent** models to compensate for errors made by previous ones.
- Or, we might train multiple models and then select a **best one** to make predictions for a given x .
- Such models are usually called **ensemble models** since they make predictions based on a **set** of trained models instead of just one.

Lecture objectives



After this lecture you will:

- Understand how **bagging** (or **bootstrap aggregation**) can be used to generate multiple datasets from a single one.
- Understand how **committees** are formed from fitting multiple models to **bootstrap datasets**.
- Understand how and why **committees** reduce model error through **averaging**.
- Understand how **boosting** works to train a sequence of **weak learners** and how these learners are **combined** into a stronger committee via weighting.
- Understand how **tree models** partition the space into **cuboid** regions and how to learn their structure and parameters.

Committees and Bagging

The power of consensus

- We saw way back in one of our first lectures how **sampling** multiple models and **averaging** helps reduce **variance**:
- Can we take advantage of this without a **fully-Bayesian** model?



Bootstrap sampling

- Consider a **regression** problem where we are predicting a single continuous target.
- The variance reduction approach worked by fitting a **low bias** model to **multiple** datasets and averaging the outputs.
- In practice we have only **one** dataset, however.
- The **bootstrap** method works by sampling M datasets of size $N' < N$ from \mathcal{D} **with replacement**.
- Each of these **bootstrap datasets** reflects the underlying distribution of \mathcal{D} but is **incomplete**.

Bootstrap aggregation (or *bagging*)

- We can then fit a model $y_m(\mathbf{x})$ to each **bootstrap dataset** (using **any** method).
- And then form the **committee** model by averaging the M **base models**:

$$y_{\text{com}}(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M y_m(\mathbf{x})$$

- Denote by $h(\mathbf{x})$ the **true** regression function generating \mathcal{D} .
- The output of each model can be written as this **true** function plus an **error**:

$$y_m(\mathbf{x}) = h(\mathbf{x}) + \varepsilon_m(\mathbf{x})$$

- Note that we haven't done **anything** yet – except identify the **error** in each of our models.

Quantifying error

- The expected **squared** error of each model is then:

$$\mathbb{E}_{\mathbf{x}} \left[\{h(\mathbf{x}) - y_m(\mathbf{x})\}^2 \right] = \mathbb{E}_{\mathbf{x}} [\varepsilon_m(\mathbf{x})^2]$$

- And the **average error** of the individual models is:

$$E_{\text{av}} = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_{\mathbf{x}} [\varepsilon_m(\mathbf{x})^2]$$

- Again**, we haven't really said anything earth shattering: the average squared error made by the models is the average of all of the squared errors made by the models.

The committee error

- We can compute the expected error of the **committee** similarly:

$$\begin{aligned} E_{\text{com}} &= \mathbb{E}_{\mathbf{x}} \left[\left\{ \frac{1}{M} \sum_{m=1}^M h(\mathbf{x}) - y_m(\mathbf{x}) \right\}^2 \right] \\ &= \mathbb{E}_{\mathbf{x}} \left[\left\{ \frac{1}{M} \sum_{m=1}^M \varepsilon_m(\mathbf{x}) \right\}^2 \right] \end{aligned}$$

- OK, but squaring that \sum is going to be a huge **mess**...

The committee error

- How can we overcome this? Answer: **assumptions!**

$$E_{\text{com}} = \mathbb{E}_{\mathbf{x}} \left[\left\{ \frac{1}{M} \sum_{m=1}^M \varepsilon_m(\mathbf{x}) \right\}^2 \right]$$

- If we assume the errors are **zero-mean** and **uncorrelated**:

$$E_{\text{com}} [\varepsilon_m(\mathbf{x})] = 0$$

$$E_{\text{com}} [\varepsilon_m(\mathbf{x}) \varepsilon_l(\mathbf{x})] = 0 \quad m \neq l$$

- Then we have:

$$E_{\text{com}} = \frac{1}{M} E_{\text{av}}$$

Averaging models reduces expected error

- So if we **average** M models we reduce the expected error of the **single** models by a factor of $1/M$!

$$E_{\text{com}} = \frac{1}{M} E_{\text{av}}$$

- Remember the **assumptions**, though: errors are almost **never** uncorrelated.
- However, there **are** guarantees that the **committee error** will never exceed the expected model error (i.e. $E_{\text{com}} \leq E_{\text{av}}$).

SKIPPABLE Boosting

SKIPPABLE Weighted averaging


- The **committee** model computes a simple average over base models:

$$y_{\text{com}}(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M y_m(\mathbf{x})$$

- These models are fit by **bagging**: multiple **samples** from training set **with replacement**.
- The base models are **independently** trained on the **bootstrap datasets** (i.e. the samples from **bagging**).
- What if we don't treat all models as **equal** in the committee?
- What if we train them in **sequence** so they "help each other out" in some way?

SKIPPABLE Boosting

- This is precisely what **boosting** does:



$$Y_M(\mathbf{x}) = \text{sign} \left(\sum_m^M \alpha_m y_m(\mathbf{x}) \right)$$

SKIPPABLE AdaBoost

- The AdaBoost algorithm does this for binary classification problems.
- Assume we have a dataset $\mathcal{D} = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\}$ for $t_n \in \{-1, 1\}$.
- The AdaBoost idea:
 - We associate with each sample a weight, which is initially $\frac{1}{N}$.
 - Assume we have a way to train a base classifier with weighted samples.
 - After training M base classifiers we combine them into a committee with coefficients giving different weights to each classifier.
- Note: in boosting, these base classifiers are usually called weak learners.

SKIPPABLE AdaBoost

AdaBoost

1. Initialize the data weighting coefficients $\{w_n\}$ by setting $w_n^{(1)} = 1/N$ for $n = 1, \dots, N$.
2. For $m = 1, \dots, M$:
 - (a) Fit a classifier $y_m(\mathbf{x})$ to the training data by minimizing the weighted error function

$$J_m = \sum_{n=1}^N w_n^{(m)} I(y_m(\mathbf{x}_n) \neq t_n) \quad (14.15)$$

where $I(y_m(\mathbf{x}_n) \neq t_n)$ is the indicator function and equals 1 when $y_m(\mathbf{x}_n) \neq t_n$ and 0 otherwise.

- (b) Evaluate the quantities

$$\epsilon_m = \frac{\sum_{n=1}^N w_n^{(m)} I(y_m(\mathbf{x}_n) \neq t_n)}{\sum_{n=1}^N w_n^{(m)}} \quad (14.16)$$

and then use these to evaluate

$$\alpha_m = \ln \left\{ \frac{1 - \epsilon_m}{\epsilon_m} \right\}. \quad (14.17)$$

- (c) Update the data weighting coefficients

$$w_n^{(m+1)} = w_n^{(m)} \exp \{ \alpha_m I(y_m(\mathbf{x}_n) \neq t_n) \} \quad (14.18)$$







3. Make predictions using the final model, which is given by

$$Y_M(\mathbf{x}) = \text{sign} \left(\sum_{m=1}^M \alpha_m y_m(\mathbf{x}) \right). \quad (14.19)$$

SKIPPABLE AdaBoost

- The ε_m represent the (weighted) error rates of the classifier y_m .
- The weighting coefficients α_m assign higher weight to more accurate classifiers in (14.19).
- And are used to give higher weight to misclassified samples in (14.18).
- AdaBoost works very well in practice, especially with many, potentially very weak learners.


SKIPPABLE AdaBoost



Tree Models

Recursive space partitioning

- We can think of **classification** (also regression) as making **constant** predictions on **partitions** of the input space:



Decision Trees

- Tree models are simple but *widely used* models that work by partitioning.
- Axis-aligned trees partition the space into **cuboids** aligned with the axes.
- They are **ensemble models**: a **single** model is responsible for each **partition**.
- To select **which** model we **traverse** the tree in a **sequential decision process**.
- Tree models are often considered **interpretable** by humans.
- The model we consider is the **Classification and Regression Tree (CART)**.

Learning a tree model

- Consider again a regression problem with dataset
 $\mathcal{D} = \{(x_1, t_1), \dots, (x_N, t_N)\}$, where now our t_n are continuous targets.
- We wish to partition the space \mathbb{R}^D so that in each partition our estimate of the target there minimizes the squared error.
- This is pretty easy: we just need the average of the targets t_n of all points x_n falling in that partition.
- The real problem is *how should we determine this partitioning*.
- Optimizing over all possible partitionings is combinatorially intractable, even for axis-aligned trees...

Greedy learning

- Instead, we will use a recursive, **greedy** policy:
 1. Start from a single **root node** corresponding to the whole space.
 2. Iterate over each of the D possible splitting variables:
 - 2.1 Consider each possible **threshold** to split the data into two sets.
 - 2.2 Select the one that **minimizes** the sum-of-squared errors in the splits.
 3. Recursively apply this procedure to the **children**.
- When should we **stop** this procedure?

Greedy learning

- Instead, we will use a recursive, **greedy** policy:
 1. Start from a single **root node** corresponding to the whole space.
 2. Iterate over each of the D possible splitting variables:
 - 2.1 Consider each possible **threshold** to split the data into two sets.
 - 2.2 Select the one that **minimizes** the sum-of-squared errors in the splits.
 3. Recursively apply this procedure to the **children**.
- When should we **stop** this procedure?
- The usual strategy is to **oversegment** the input space by building a **deep and broad** tree.
- Why might this be problematic?

Greedy learning

- Instead, we will use a recursive, **greedy** policy:
 1. Start from a single **root node** corresponding to the whole space.
 2. Iterate over each of the D possible splitting variables:
 - 2.1 Consider each possible **threshold** to split the data into two sets.
 - 2.2 Select the one that **minimizes** the sum-of-squared errors in the splits.
 3. Recursively apply this procedure to the **children**.
- When should we **stop** this procedure?
- The usual strategy is to **oversegment** the input space by building a **deep and broad** tree.
- Why might this be problematic?
- And then **prune back** the tree to balance **complexity** and **error minimization** on the training data.

Greedy learning

- Assume we already have a partitioning (a **tree**), with leaf nodes indexed by $\tau = 1, \dots, T$.
- Call the **partition** corresponding to each leaf \mathcal{R}_τ .
- The **optimal prediction** and **squared error** in at each node are:

$$y_\tau = \frac{1}{N_r} \sum_{x_n \in \mathcal{R}_\tau} t_n \quad Q_\tau(T) = \sum_{x_n \in \mathcal{R}_\tau} (t_n - y_\tau)^2$$

- So we can define a **pruning criterion** like:

$$C(T) = \sum_{\tau=1}^{|T|} Q_\tau(T) + \lambda |T|$$

For classification problems

- If we have a **classification** problem, we just need to change the **error**.
- Define $p_{\tau k}$ as the proportion of samples of class k at node τ .
- For example the **negative cross entropy**:

$$Q_\tau(T) = \sum_{k=1}^K p_{\tau k} \ln p_{\tau k}$$

- Or the **Gini Index**:

$$Q_\tau(T) = \sum_{k=1}^K p_{\tau k} (1 - p_{\tau k})$$

Concluding Remarks

The power of consensus (and diversity!)

- Ensemble models are a powerful tool for combining simple models into better ones.
- Committees reduce the average error rate – but only if the assumptions are valid.
- Trees are interpretable models that partition the input space into homogeneous regions.

Reading and Homework Assignments

Reading Assignment:

- Bishop: Chapter 14 (14.1, 14.2, 14.3, 14.4, 14.5)
- Scikit-learn User Guide (chapter 1.11)