

Fundamentals of Machine Learning:

Introduction to Deep Learning

Prof. Andrew D. Bagdanov ([andrew.bagdanov AT unifi.it](mailto:andrew.bagdanov@unifi.it))



UNIVERSITÀ
DEGLI STUDI
FIRENZE

DINFO
DIPARTIMENTO DI
INGEGNERIA
DELL'INFORMAZIONE

Outline

Introduction

Connectionism: The Old School

Connectionism: The New School

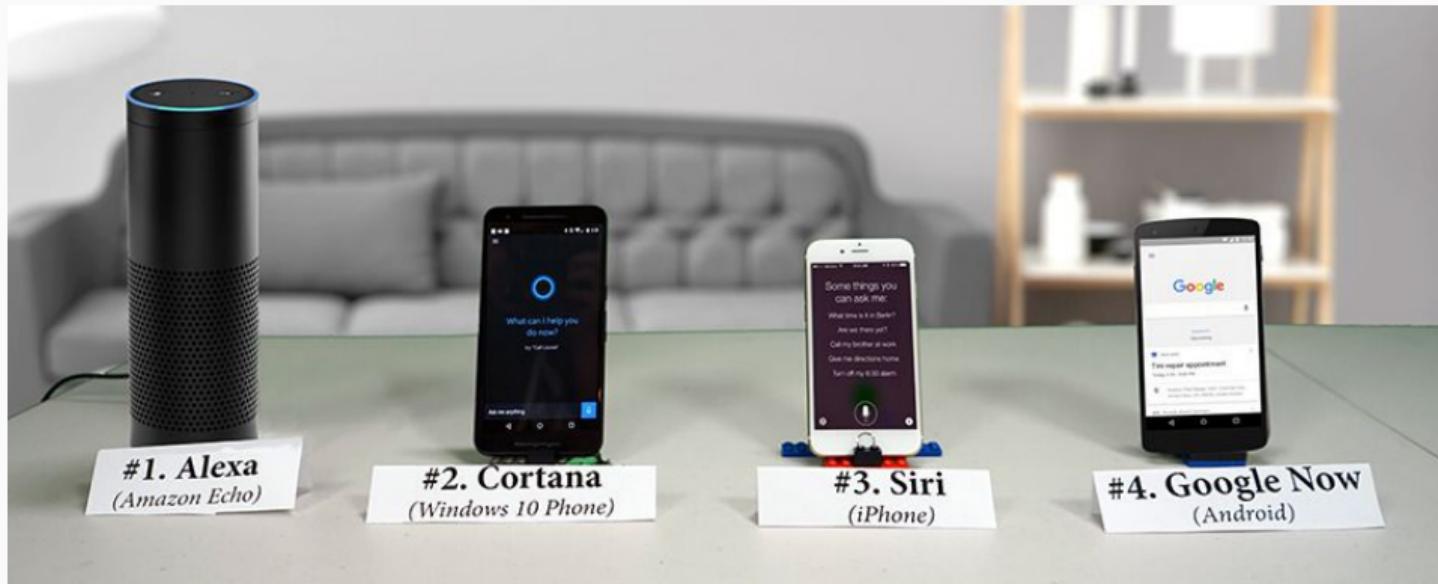
Building and Training Deep Networks

Discussion

Introduction

Digital assistants

- Deep learning is **profoundly** changing our lives.



Natural language processing

- Deep Recurrent Neural Networks (RNNs) are powering the latest generation of natural language translation technologies.

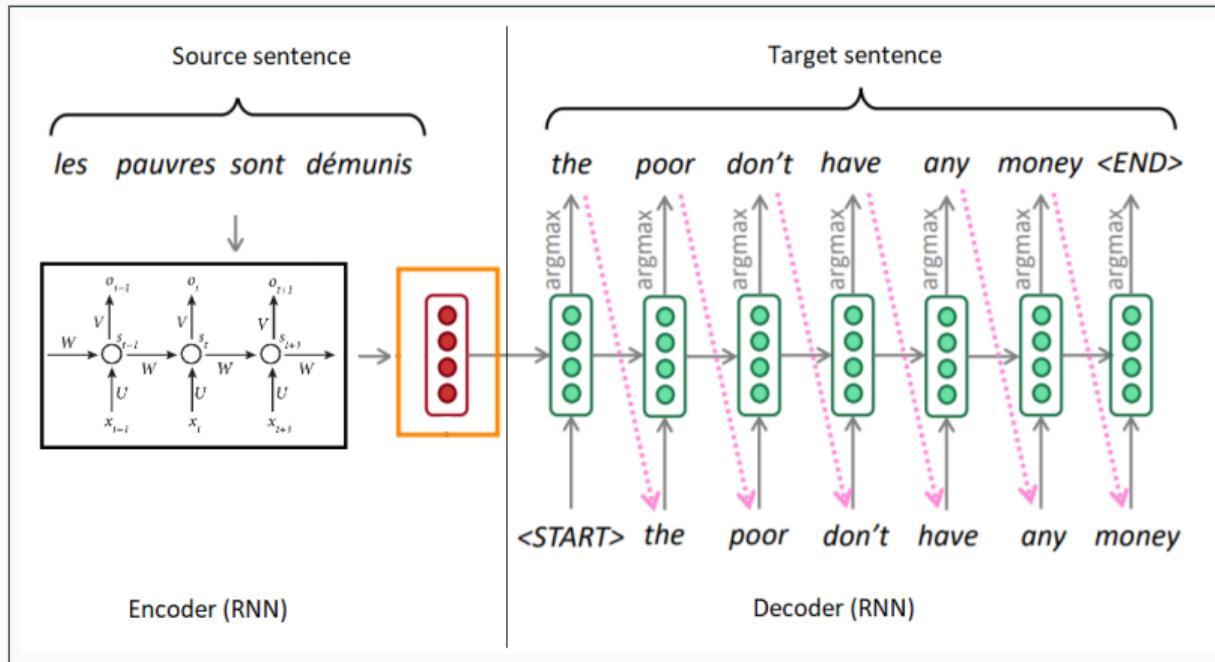


Image captioning

- Convolutional Neural Networks (CNNs) are able to extract high-level semantics from images.

Train

COCO Captions: 80 Classes



Two pug **dogs** sitting on a **bench** at the beach.



A **child** is sitting on a **couch** and holding an **umbrella**.

Open Images: 600 Classes



Goat



Artichoke



Accordion



Dolphin



Waffle



Balloon

nocaps Val / Test

In-Domain: Only COCO Classes



The **person** in the brown suit is directing a **dog**.

Near-Domain: COCO & Novel Classes



A **person** holding a black **umbrella** and an **accordion**.

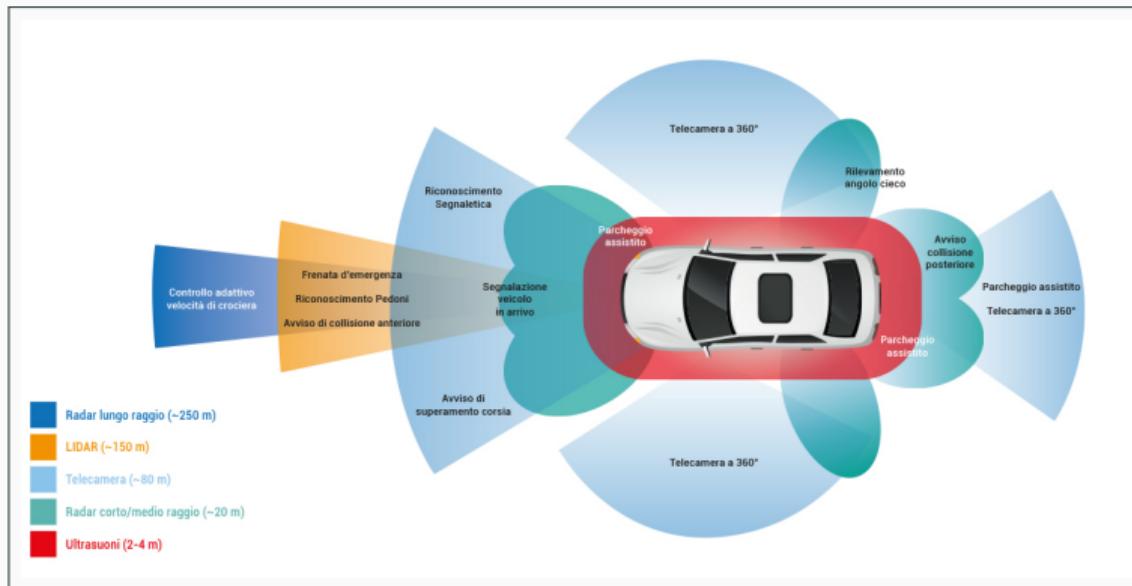
Out-of-Domain: Only Novel Classes



Some **dolphins** are swimming close to the base of the ocean.

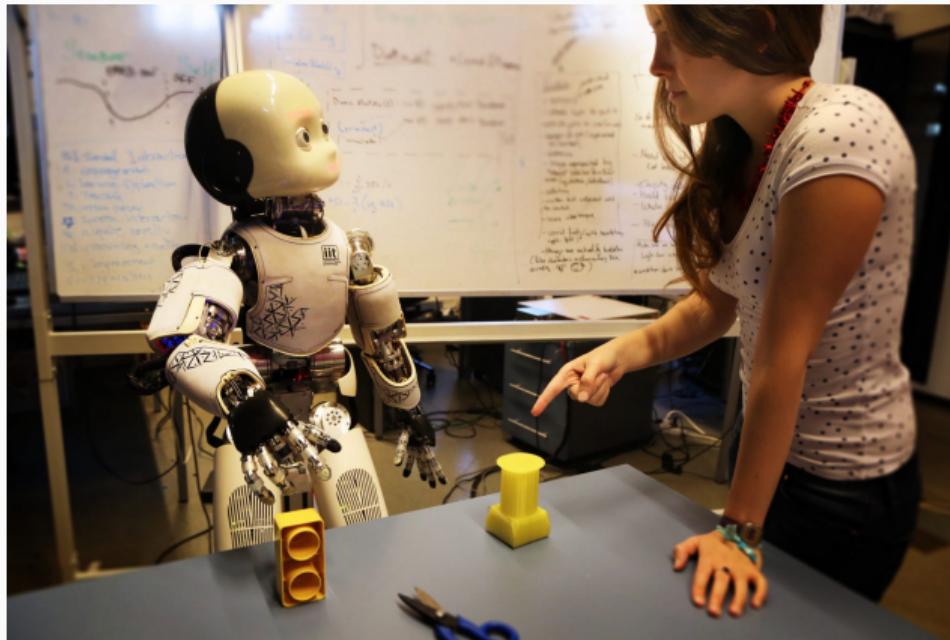
Self-driving cars

- CNNs are able to integrate multi-modal inputs and are driving the latest advances in Automatic Driving Assistance (ADAS) systems.



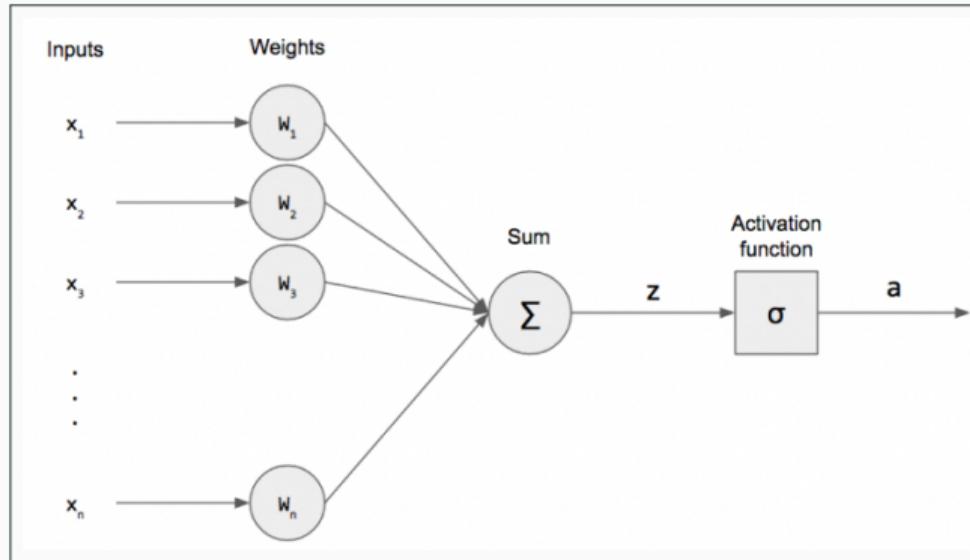
Reinforcement learning

- Deep Reinforcement Learning is being used to train robots who can learn from **experience** and **interactions** with their environment.



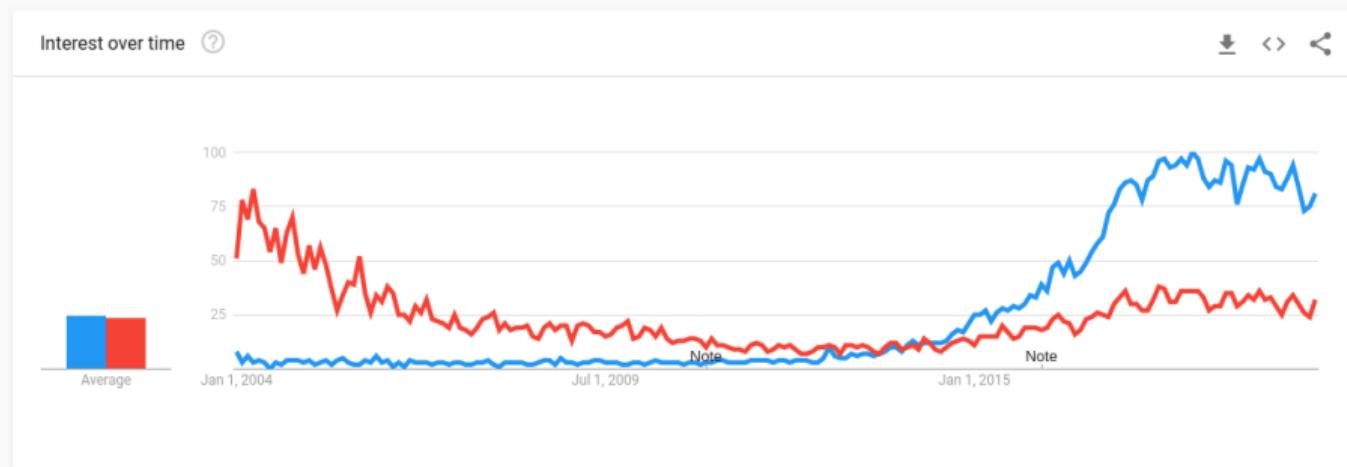
All thanks to...

- The humble Neural Network.
- Artificial Neural Networks (ANNs) are extremely simple, yet also extremely powerful models.
- They are, in fact, universal function approximators.



Neural Networks are not new

- As we will see, **neural networks** have a storied history.
- **Deep Learning**, however, is their modern incarnation.



Overview

- Today we will see what puts the **deep** into **Deep Learning**.
- We will start with an overview of some **historical milestones** in the development of artificial neural networks.
- Then we will look at how modern deep neural networks are **actually built**:
- We will see how the basic **Multilayer Perceptron (MLP)** model provides a **modular** architecture for machine learning problems.
- And we will see how modern tools (e.g. **PyTorch**) makes it easy to apply Deep Models to new problems.

Lecture objectives

After this lecture you will:

- Understand what **connectionist models** are and how they allow us to compute nonlinear functions of inputs.
- Understand how the **perceptron** works and how its parameters are estimated via **error correction**.
- Understand how the **multi-layer perceptron** generalized the perceptron and how it uses **input**, **hidden**, and **output** layers to represent feedforward computations of inputs to outputs.

Connectionism: The Old School

What is a Deep Neural Network?

- Neural Networks are connectionist models.
- Connectionism has deep roots reaching back to Classical Greece.
- To understand this rich inheritance it is useful to go back in time and trace the roots of modern Deep Models.
- Connectionism arose from the neuroscience and psychological research communities of the 1940s and 1950s.
- These were the nascent beginning of what would become Cognitive Science.
- Though founded on solid experimental practice, what was lacking was any sort of computation basis for learning.

Connectionism: Hebbian Learning

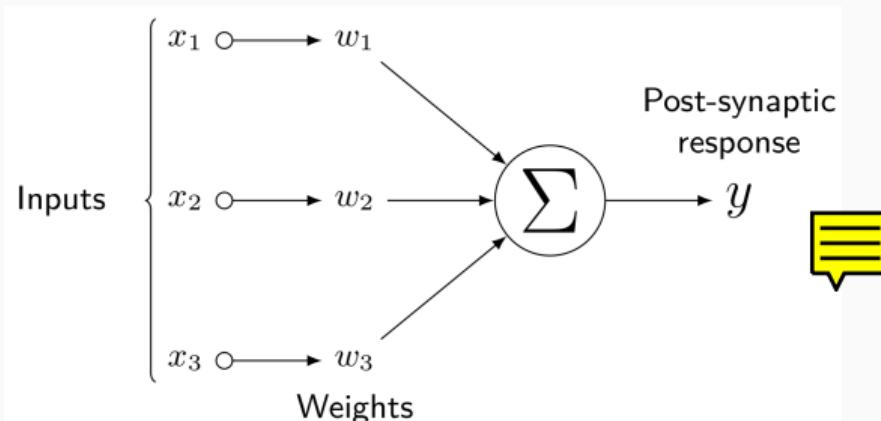
- One of the first **concrete** learning rules for connectionist models (both artificial and biological).
- **Hebb's Rule**: if cell A consistently contributes to the activity of cell B, then the synapse from A to B should be strengthened.
- **More quaintly**: *neurons that fire together, wire together; neurons that fire out of sync, fail to link.*

$$w_{ij} = x_i x_j$$

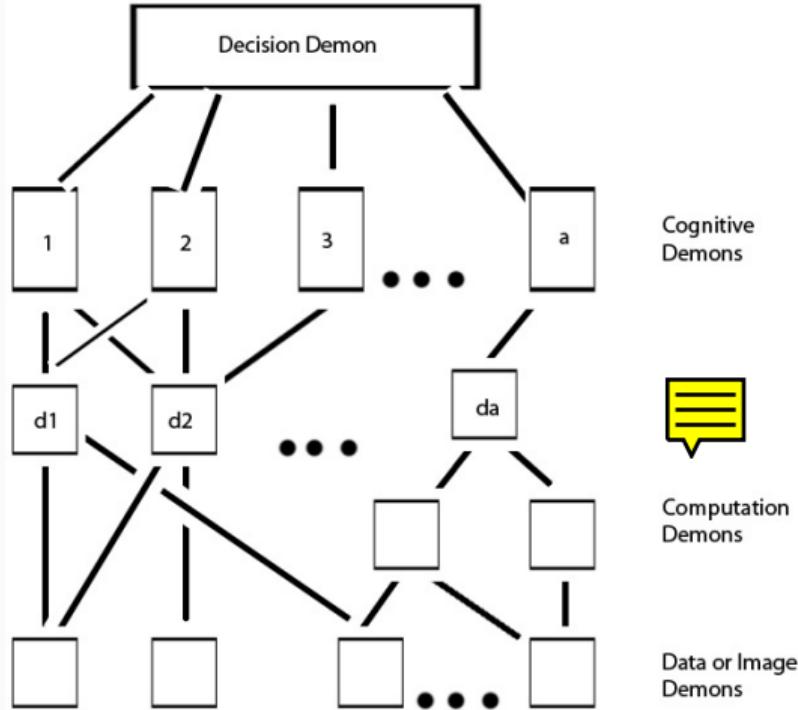
$$w_{ij} = \frac{1}{N} \sum_p x_i^p x_j^p$$

$$\Delta w_i = \eta x_i y$$

$$= \eta x_i \sum_j w_i x_i$$



Connectionism: The Pandemonium Model



- In 1958 Selfridge proposed a multi-layer, **parallel** model of machine learning.
- The model consists of four layers, each inhabited by **demons**.
- Network architecture fixed *a priori*, connections updated using **supervised** learning.
- Demons **yell upwards**, higher-level ones listen and respond.
- **High-worth** demons can replace low-worth ones via **combination**.

Connectionism: The Perceptron



- The **Perceptron** is probably the simplest (and most famous) feedforward neural network.
- The **perceptron algorithm** was invented by Rosenblatt in 1958.
- It was designed to be a **machine**, and its original purpose was to perform **image recognition**.

$$f(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$



The perceptron algorithm

Input: $D = \{(x_i, y_i)\}_{i=1}^N$ (training data)

Output: learned weights w

$w_0 \leftarrow$ random initialization

$t \leftarrow 1$

while not converged **do**

for $(x, y) \in D$ **do**

$$\hat{y} = f(w^T x)$$



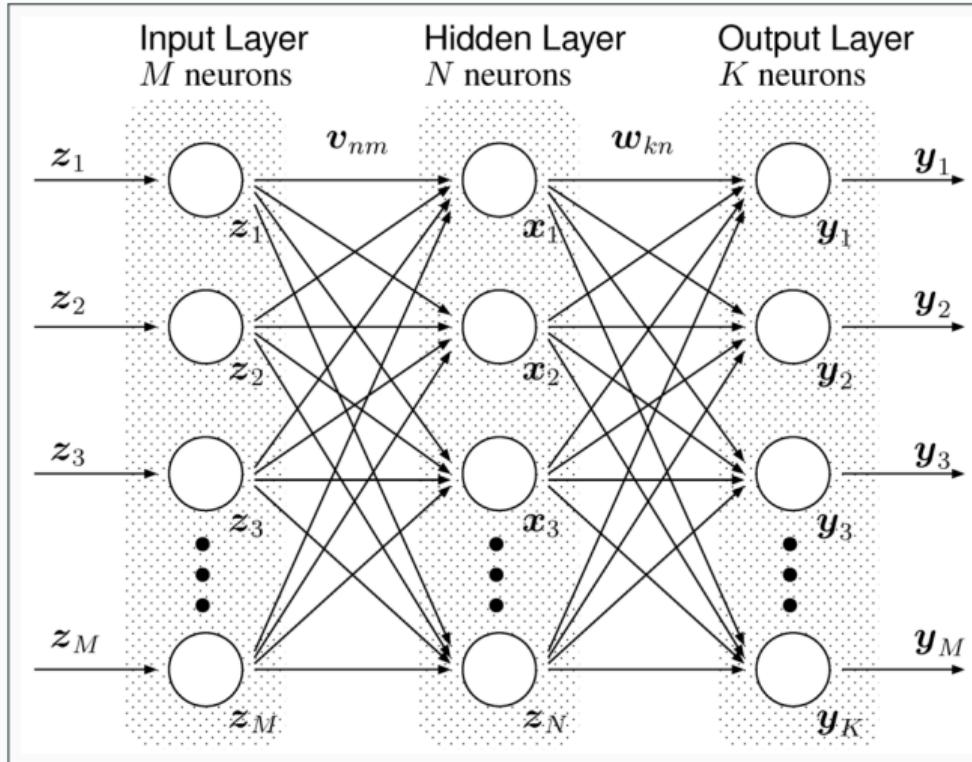
$$w_t \leftarrow w_{t-1} + \eta(y - \hat{y})x$$

$$t \leftarrow t + 1$$

Connectionism: The New School

The Multilayer Perceptron

- Let's look at a simple Neural Network architecture known as the Multilayer Perceptron (MLP):



The Multilayer Perceptron (continued)

- The MLP equation (one hidden layer):

$$\hat{y}(x) = \sigma(w_2^T \sigma(w_1^T x + b_1) + b_2)$$



- Except for the activation function σ , this is a linear system.
- Common activation functions (elementwise):

- $\sigma(x) = \tanh(x)$
- $\sigma(x) = (1 + e^{-x})^{-1}$
- $\sigma(x) = \frac{\exp(x)}{\sum_i \exp(x_i)}$ (softmax, used for outputs).



Training an MLP

- How do you train a model?
- Decide on a loss function (like the negative log-likelihood):

$$L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x})) = -\frac{1}{C} \sum_i y_i \log(\hat{y}_i)$$

- And perform gradient descent w.r.t. all model parameters:

$$\boldsymbol{\theta}_{n+1} = \boldsymbol{\theta}_n - \varepsilon \nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}))$$

$$\boldsymbol{\theta}_{n+1} = \boldsymbol{\theta}_n - \varepsilon \frac{1}{N} \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}_i))$$



- Where ε is the learning rate.
- The standard algorithm for this is known as backpropagation and it is very clever and efficient.

Training an MLP (continued)

- OK, we perform **gradient descent** like this:

$$\boldsymbol{\theta}_{n+1} = \boldsymbol{\theta}_n - \varepsilon \frac{1}{N} \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}_i))$$

Training an MLP (continued)

- OK, we perform **gradient descent** like this:

$$\boldsymbol{\theta}_{n+1} = \boldsymbol{\theta}_n - \varepsilon \frac{1}{N} \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}_i))$$

- But... How do I compute that gradient $\nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}_i))$ for **non-trivial** models?

Training an MLP (continued)

- OK, we perform gradient descent like this:

$$\boldsymbol{\theta}_{n+1} = \boldsymbol{\theta}_n - \varepsilon \frac{1}{N} \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}_i))$$

- But... How do I compute that gradient $\nabla_{\boldsymbol{\theta}} L(\mathbf{y}, \hat{\mathbf{y}}(\mathbf{x}_i))$ for non-trivial models?
- The standard algorithm for this is known as backpropagation and it is very clever and efficient.
- But... We will defer a detailed discussion of backpropagation for the next lecture.

Stochastic Gradient Descent (SGD)

- Problem: what happens if N (the number of training samples) is very large?
- Well, we end up taking very slow steps – each iteration of gradient descent is an average over the entire dataset.
- Solution: approximate the true gradient with the gradient at a single training example:

Online Stochastic Gradient Descent



- Choose an initial vector of parameters θ and learning rate η .
- Repeat until an approximate minimum is found:
 1. Randomly shuffle training samples in D .
 2. For $(x, y) \in D$:
$$\theta := \theta - \eta \nabla_{\theta} \mathcal{L}(\{x, y\}; \theta)$$

Stochastic Gradient Descent (continued)

- Another problem: evaluating the gradient on single examples leads to very noisy steps in parameter space.
- One trick to mitigate this is to use momentum: keep a running average of gradients that is slowly updated.
- Another solution is to use mini-batches: instead of a single sample, average the gradients over a small batch of samples. 
- It is common to use a combination of mini-batches and momentum to stabilize training.

SGD Terminology

- Some useful terminology for deep learning optimization:

- 1 epoch: one complete pass over the data.
- 1 iteration: a single gradient step.
- N : number of training samples.
- B : batch size.



Algorithm	iterations per epoch
Batch gradient descent	1
Stochastic Gradient Descent	N
Mini-batch Gradient Descent	$\frac{N}{B}$

Building and Training Deep Networks

The ingredients

- **Network Definition:** A neural network **model** must be defined that precisely describes the transformation from input to output (e.g. a Multi-layer Perceptron with one hidden layer). The model is typically – but not necessarily – defined in terms of pre-defined, composable **modules**.
- **Dataset and DataLoader:** To train a model, we need some **data** (duh). Data management is **critical** in Deep Learning, and a **DataLoader** is responsible for loading, transforming, and **batching** data for training and inference.
- **A Training Loop:** Neural Networks are not **black boxes**, and their training can be **delicate** and **subtle**. A **training loop** applies an iterative optimization algorithm over **training batches**.



Defining a network: Basics

- A **network architecture** defines the family of functions we use as **models**.
- It explicitly defines the **parameterized** family of functions \mathcal{H} we are searching over.
- An architecture is a **Directed Acyclic Graph (DAG)** defining the connections between basic **computational blocks**.
- We call these blocks **layers** and they are reusable **building blocks**.

Defining a network: The PyTorch Tensor

- Tensors are a specialized data structure similar to arrays and matrices.
- In PyTorch, tensors are used to:
 - Encode the inputs and outputs of a model.
 - Encode the parameters of a model.
- Tensors vs. NumPy ndarrays:
 - Tensors can run on GPUs or other hardware accelerators.
 - Tensors and NumPy arrays can share the same underlying memory, eliminating data copying.
- Tensors are optimized for automatic differentiation.



Defining a network: Tensor Operations

- PyTorch tensors support over **100 tensor operations** for:
 - Arithmetic
 - Linear algebra
 - Matrix manipulation (transposing, indexing, slicing)
 - Sampling
 - And much, much more.
- Operations can be run on **GPU** (typically faster than CPU)
 - Default tensor creation on CPU
 - Move to GPU using `.to` method (check GPU availability first)

Defining a network: The tedious way

- We can define our network model completely using low-level tensors.
- For linear regression it might look something like this:

```
# Define our model parameters, initialize randomly.  
w = torch.randn(Xs.shape[1])  
b = torch.randn(1)  
  
# Try it out by predicting on training set.  
Xs_tr @ w + b  
  
--> tensor([-428.8886, -384.7284, -212.2810, ..., -286.6152, -326.9498, -162.2029])
```

Defining a network: The tedious way

- OK, but that just looks like NumPy **with extra steps**.
- The real magic is in the construction of the **computational graph**.
- And what we can **do** with the computational graph.

```
# A quadratic loss function.
def l2loss(y, y_hat):
    return ((y - y_hat) ** 2.0).mean()

# Define our model parameters, initialize randomly.
w = torch.randn(Xs.shape[1], requires_grad=True)
b = torch.randn(1, requires_grad=True)

# Compute the loss on the training set.
l2loss(ys_tr, Xs_tr @ w + b)

--> tensor(909343.6875, grad_fn=<MeanBackward0>)
```

Defining a network: The tedious way

- Our computational graph is **instrumented** to compute gradients!
- Gradients of the loss wrt tensors with **requires_grad=True** will be computed:

```
# Compute the gradient of the loss function.  
loss = l2loss(ys_tr, Xs_tr @ w + b)  
loss.backward()  
print(f'Gradient wrt w: {w.grad}')  
print(f'Gradient wrt b: {b.grad}')  
  
-->  
Gradient wrt w: tensor([-6.7520e+03, -4.5452e+04, -9.1817e+03, -1.8607e+03,  
                         -4.0045e+06, -6.9095e+03, -6.1650e+04,  2.0769e+05])  
Gradient wrt b: tensor([-1739.1476])
```

Defining a network: The tedious way

- We **could** perform gradient descent, but everything must be done **just right**:

```
# Define our model parameters, initialize randomly.  
w = torch.normal(0.0, np.sqrt(1.0/((8+1)/2)), size=(Xs.shape[1],)).requires_grad_()  
b = torch.tensor(1.0, requires_grad=True)  
  
# Standardize dataset.  
Xs_tr -= Xs_tr.mean(0)  
Xs_tr /= Xs_tr.std(0)
```

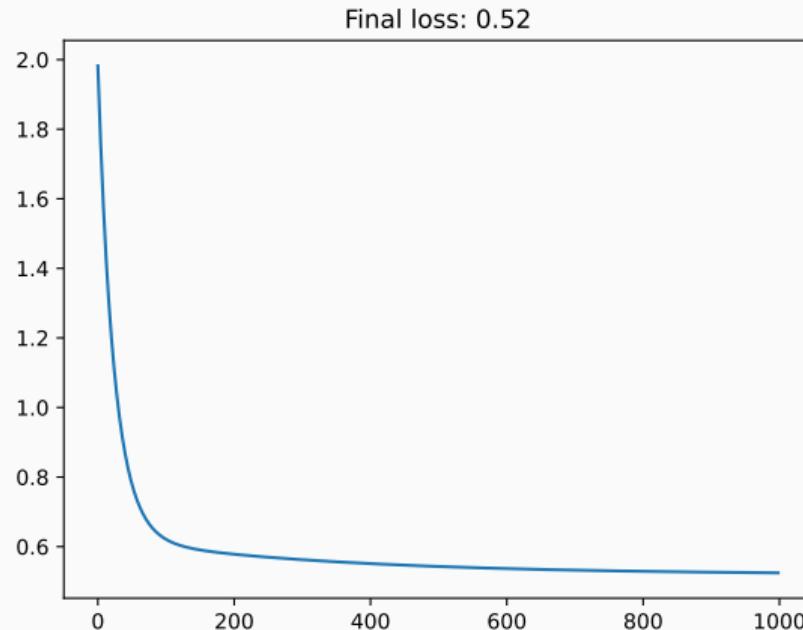
Defining a network: The tedious way

- We could perform gradient descent, but everything must be done just right:

```
# Training hyperparameters.  
epochs = 1000  
lr = 1e-2  
  
# Training loop.  
losses = []  
for it in range(epochs):  
    loss = l2loss(ys_tr, Xs_tr @ w + b)  
    loss.backward()  
    w.data -= lr * w.grad.data  
    b.data -= lr * b.grad.data  
    w.grad.data.zero_()  
    b.grad.data.zero_()  
    losses.append(loss.item())  
  
# Plot losses and report last loss (MSE).  
plt.plot(losses[1:]); plt.title(f'Final loss: {losses[-1]}')
```

Defining a network: The tedious way

- We **could** perform gradient descent, but everything must be done **just right**



Defining a network: The modular way

- OK, that **really** sucks. And is **error-prone**. And doesn't **scale**. And has lots of **other problems**.
- Instead, we will make use of the **reusable modules** from the **torch.nn** package.
- This **PyTorch namespace** contains implementations of **very many** types of layers. Let's take a look...
- The simplest way to define our **linear regression model** is:

```
model = nn.Linear(8, 1)
print(list(model.parameters()))
--> [Parameter containing:
    tensor([-0.1241,  0.2113, -0.1457, -0.2680, -0.3201, -0.2886, -0.2482, -0.0639]),
        requires_grad=True),
Parameter containing: tensor([0.3110], requires_grad=True)]
```

Defining a network: The modular way

- The classes from the `torch.nn` package **encapsulate** all parameters and provide a uniform **API** to access them.
- Their **instances** act like **functions**:

```
# Our linear model.  
model = nn.Linear(8, 1)  
  
# Training loop.  
losses = []  
for it in range(epochs):  
    loss = l2loss(ys_tr, model(Xs_tr).flatten())  
    model.zero_grad()  
    loss.backward()  
    for p in model.parameters():  
        p.data -= lr * p.grad.data  
    losses.append(loss.item())
```

Defining a network: The modular way

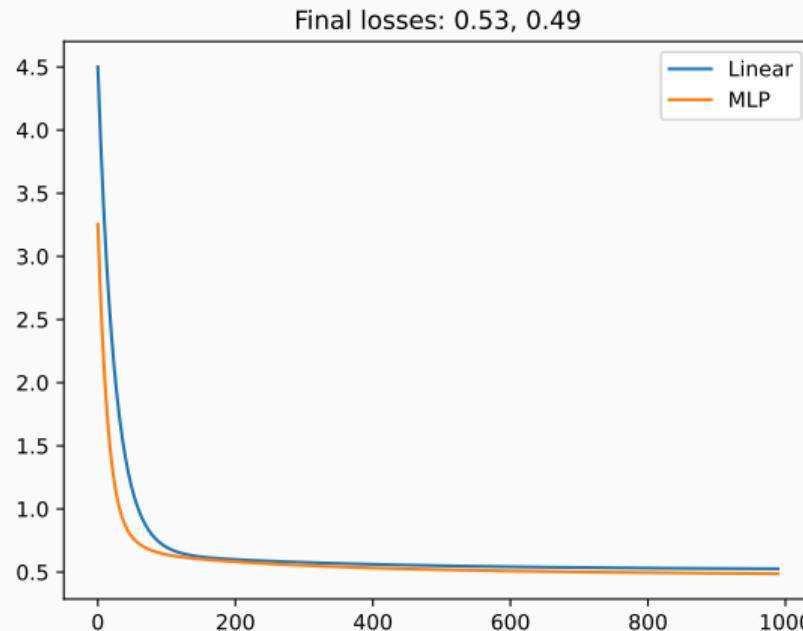
- The **modularity** really shines when layers are **composed**:

```
# Our MLP model.  
model = nn.Sequential(nn.Linear(8, 16), nn.Tanh(), nn.Linear(16, 1))  
  
# Training loop.  
losses_mod = []  
for it in range(epochs):  
    loss = l2loss(ys_tr, model(xs_tr).flatten())  
    model.zero_grad()  
    loss.backward()  
    for p in model.parameters():  
        p.data -= lr * p.grad.data  
    losses_mod.append(loss.item())
```



Defining a network: The modular way

- Which makes **model selection** reasonably agile:



Defining a network: The modular way

- The most general way to define models is to extend `nn.Module`.
- This way we have complete access to everything and control over composition:

```
class LinRegMLP(nn.Module):  
    def __init__(self, num_in, num_hidden, activation=F.tanh):  
        super().__init__()  
        self.activation = activation  
        self.fc1 = nn.Linear(num_in, num_hidden)  
        self.out = nn.Linear(num_hidden, 1)  
  
    def forward(self, xs):  
        return self.out(self.activation(self.fc1(xs)))
```

Dataset and DataLoader: Basics

- Data is at the **heart** of Deep Learning.
- The important breakthroughs in Deep Learning would not have been possible without **massive** datasets.
- **Annotated** or **unlabeled** data may power training of Deep Models.
- Model architectures are **strongly** related to the **type** of data.
- Dataset **size** and **annotation quality** will determine how large of a model we can train.

Dataset and DataLoader: Batching

- In the examples above we are performing **batch gradient descent**.
- Usually this is **inadvisable**, so we use a **DataLoader** which is a **batched iterator** over a dataset:

```
from torch.utils.data import DataLoader, TensorDataset

# Create a TensorDataset
ds_train = TensorDataset(Xs_tr, ys_tr)

# Create a DataLoader
dl_train = DataLoader(ds_train, batch_size=512, shuffle=True)

for (Xs, ys) in dl_train:
    print(f'Xs: {Xs.shape}, ys: {ys.shape}')
-->
Xs: torch.Size([512, 8]), ys: torch.Size([512])
Xs: torch.Size([512, 8]), ys: torch.Size([512])
...
Xs: torch.Size([120, 8]), ys: torch.Size([120])
```



Dataset and DataLoader: The rest of the story

- Datasets for Deep Learning tend to be **massive**, and data loading is a **critical** issue.
- The **DataLoader** might be responsible for **loading** data from disk, **transforming it** (e.g. **standardization** or **data augmentation**), and **batching** it for training and validation.

Optimization: Training Deep Models

- Without optimization we cannot get anything out of data.
- There are a variety of optimization algorithms that manage training of Deep networks.
- 99% of them are based on gradient-based optimization (i.e. Stochastic Gradient Descent).
- There is also a (rather large) bag of tricks needed to avoid common training pitfalls.

Optimization: Training Deep Models

- The `torch.nn` module provides classes to **encapsulate** details of model definitions.
- In the `torch.optim` package you will find implementations of **several** optimization algorithms the **encapsulate** the details of model training:
 - `torch.optim.SGD`: Implements the standard **Stochastic Gradient Descent (SGD)** algorithm, with optional **momentum**.
 - `torch.optim.Adam`: Implements the **Adaptive Moment Estimation (Adam)** algorithm, which augments SGD with adaptive, per-parameter learning rates.

Optimization: Putting Everything Together

- First some imports and a model definition:

```
import torch.nn as nn
import torch.nn.functional as F
from torch.optim import SGD, Adam
from tqdm.notebook import tqdm

# Define our model.
class LinRegMLP(nn.Module):
    def __init__(self, num_in, num_hidden, activation=F.tanh):
        super().__init__()
        self.activation = activation
        self.fc1 = nn.Linear(num_in, num_hidden)
        self.out = nn.Linear(num_hidden, 1)

    def forward(self, xs):
        return self.out(self.activation(self.fc1(xs)))
```

Optimization: Putting Everything Together

- And a function to **train** the model:

```
def train(model, dl_train, dl_val=None, epochs=100, lr=1e-3):
    opt = Adam(model.parameters(), lr=lr) # We'll use Adam.
    (train_losses, val_losses) = ([], [])

    # Iterate for required number of epochs.
    for epoch in tqdm(range(epochs)):
        # Iterate over all batches.
        epoch_loss = 0.0
        for (Xs, ys) in dl_train:
            ys_hat = model(Xs).flatten() # Forward pass and loss computation.
            loss = F.mse_loss(ys, ys_hat)
            opt.zero_grad()           # Take a gradient step.
            loss.backward()
            opt.step()
            epoch_loss += loss.item()
        train_losses.append(epoch_loss / len(dl_train)) # Collect average epoch loss.

        # If we have a validation dataloader, compute average loss over it.
        if dl_val:
            with torch.no_grad():
                val_losses.append(np.mean([F.mse_loss(ys, model(Xs).flatten()).item() for (Xs, ys) in dl_val]))
    return (train_losses, val_losses)
```

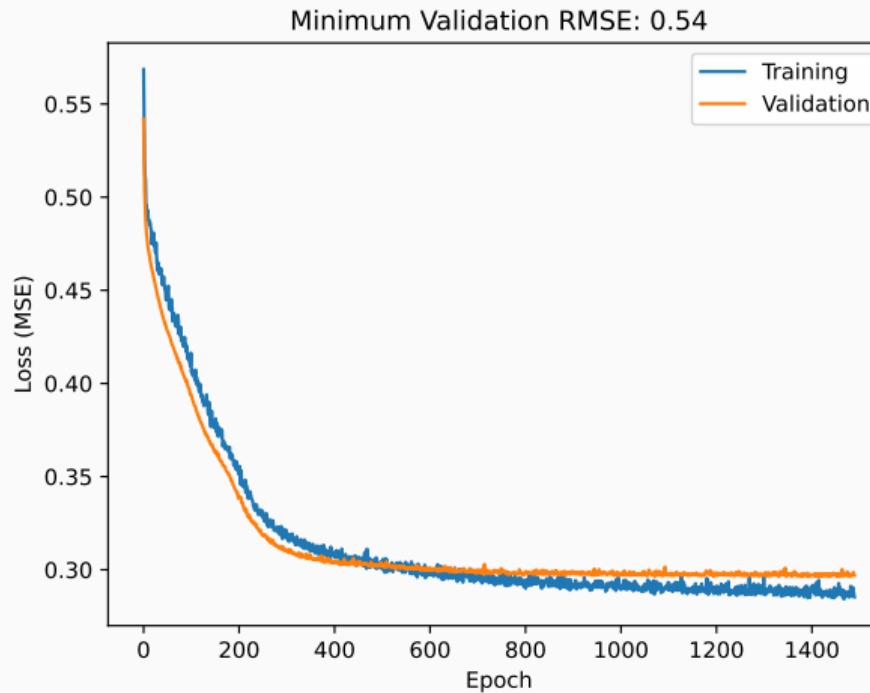
Optimization: Putting Everything Together

- And (finally) we can **use** this training loop:

```
# Training hyperparameters.  
lr = 1e-3  
epochs = 1500  
batch_size = 512  
inner_size = 16  
  
# Create datasets and DataLoaders.  
ds_train = TensorDataset(Xs_tr, ys_tr)   
ds_val = TensorDataset(Xs_te, ys_te)  
dl_train = DataLoader(ds_train, batch_size=batch_size, shuffle=True)   
dl_val = DataLoader(ds_val, batch_size=batch_size, shuffle=False)  
  
# Instantiate model, train it, and plot losses.  
model = LinRegMLP(8, inner_size)  
(train_losses, val_losses) = train(model, dl_train, dl_val, lr=lr, epochs=epochs)
```

Optimization: Putting Everything Together

- And arrive at something like this:



Discussion

Connectionist models and the state-of-the-art

- Deep Neural Networks are connectionist models that represent the state-of-the-art for many learning problems today.
- They are usually trained with backpropagation.
- Since they typically have *very many parameters* they must be trained with extreme care.
- They also typically require more training data than classical approaches in order to generalize.
- Note that there is nothing fundamentally different about these models with respect to the classical models we have already seen.
- They simply represent a much richer hypothesis set of functions \mathcal{H} over which we search.

The Embarrassment of Choice

- PyTorch is the **de facto** framework of choice for most Deep Learning projects.
- It is not the **only one**:
 - Tensorflow/Keras: a graph-based, automatic differentiating framework for deep learning in Python – now with JAX and PyTorch backends!
 - JAX: A Python library for accelerator-oriented array computation and program transformation – sold as **differentiable Python programming**.
- Moreover, there are many **frameworks** for tracking experiments:
 - Tensorboard: A suite of web applications for inspecting and understanding your TensorFlow (and PyTorch, and JAX, and Keras, and...) runs and graphs.
 - Weights and Biases: Track, version and visualize with just 5 lines of code, reproduce any model checkpoints, monitor CPU and GPU usage in real time.
 - Comet: Provides an end-to-end model evaluation platform for AI developers, with best in class LLM evaluations, experiment tracking, and production monitoring.

Reading and Homework Assignments

Reading Assignment:

- Bishop: Chapter 5
- PyTorch: [This is an excellent introductory tutorial](#)