

# Fundamentals of Machine Learning:

## Self-attention Networks and Course Wrapup

---

Prof. Andrew D. Bagdanov ([andrew.bagdanov AT unifi.it](mailto:andrew.bagdanov@unifi.it))



UNIVERSITÀ  
DEGLI STUDI  
**FIRENZE**

**DINFO**  
DIPARTIMENTO DI  
INGEGNERIA  
DELL'INFORMAZIONE

# Outline

Introduction

Self-attention Networks

From Self-attention to the Transformer

Zero- and Few-shot **Contextual** Learning

GPT-X is **Not** ChatGPT

Discussion

## Introduction

---

## Lecture Objectives

- This lecture aims **only** to give a **broad** overview of **self-attention layers** and the **Transformer** network architecture.
- None of this will be on the exam.
- Relax and enjoy this look at the **current state-of-the-art**.

## Self-attention Networks

---

## Variable length inputs and outputs

- Many types of input cannot be **easily** modeled as **vectors of fixed dimensionality** (e.g.  $\mathbb{R}^d$ ).
- Similarly, some **outputs** might not be easily modeled as vectors of fixed dimensionality.
- A classic example is **machine translation**:

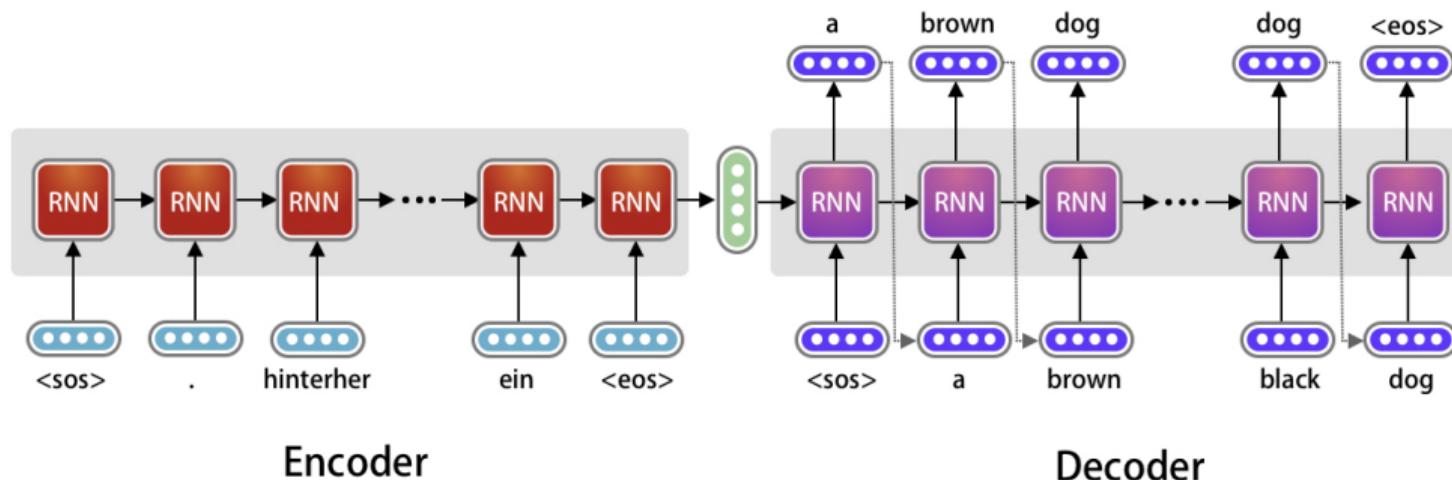
The black cat is on the wooden table.

--> Il gatto nero è sul tavolo di legno.

I wonder what Santa Claus will bring me for Christmas?

--> Mi chiedo cosa mi porterà Babbo Natale per Natale?

# The classical approach: Recurrent Neural Networks (RNNs)



- This **was** the state-of-the-art, but has **many** problems.

## The New School: Input embedding (Tokenization)

- First, we need to encode the **input sequence** into **tokens**.
- This involves learning a mapping from a sequence of **one-hot** vectors into vectors in a **continuous vector space**.
- Let  $S = [\mathbf{w}_1^T; \mathbf{w}_2^T; \mathbf{w}_3^T; \dots; \mathbf{w}_n^T]$  be a matrix whose **rows** are **one-hot** vectors in  $R^D$  ( $D$  is the size of the **vocabulary** and can be **very large**).
- We can then **embed** these words into a new space like by multiplying it by an **embedding matrix**:

$$T_0 = SW_e$$

- If  $W_e \in \mathbb{R}^{D \times d}$  the embedding space has dimensionality  $d$  (typically  $d \ll D$ ).
- $W_e$  can be **learned** or we can use a standard **tokenizer** (e.g. from **HuggingFace**).

## The New School: Queries, Keys, and Values

- Let's think for a minute about how old-school **image search engines** worked.
- A **self-attention layer** starts by mapping input tokens into **three** independent representations.
- This is done using our old friend the **linear layer** (without bias):

$$Q = T_0 W_q \quad K = T_0 W_k \quad V = T_0 W_v$$

- Our **queries**  $Q$  will be **compared** to **keys**  $K$  and the resulting **similarities** used to combine **values**  $V$ .
- This will be done for **all pairs of input tokens**.

## The New School: Self-attention

- The purpose of attention is: for each output in the sequence, predict which input tokens to focus on and how much.
- We compare queries and keys using inner products (i.e. cosine similarity).
- But, we want our combination of values to be an affine combination (coefficients sum to 1).
- So, our attention weights are computed as:

$$A = \text{softmax}(QK^T) \text{ (softmax works along rows).}$$

- And the values are combined to form each output token:

$$T_1 = AV.$$

- And we have transformed our input into new tokens.

## The New School: A simple example

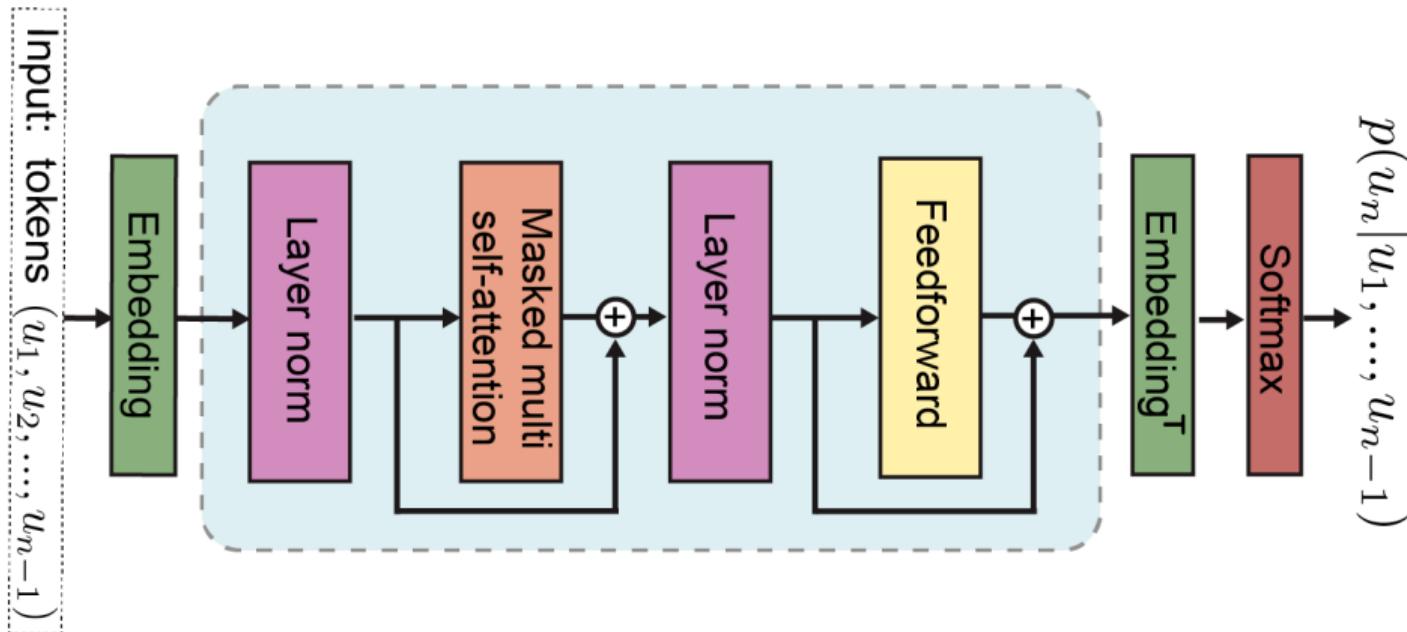
- First, a deep breath.

## From Self-attention to the Transformer

---

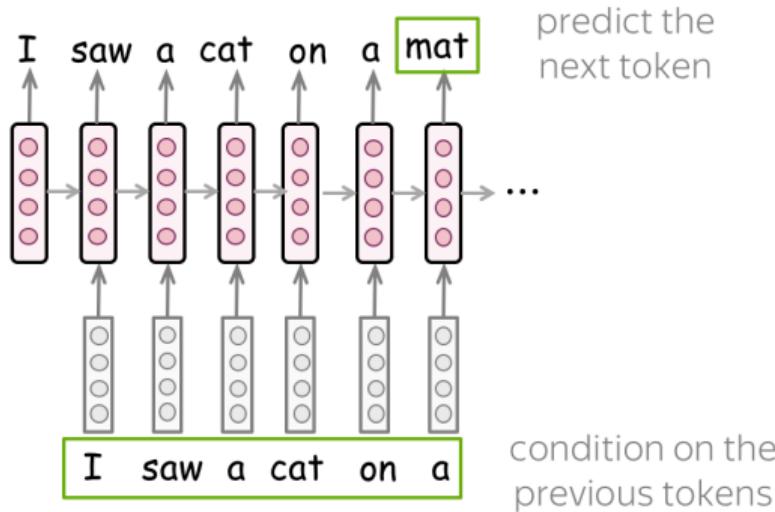
# Autoregressive training

- To build a Transformer we **stack** multiple transformer layers in sequence.
- As usual, there are a **lot** of extra details.



# Autoregressive training

- But that is **basically all there is**.
- GPT-2/3 were trained **exclusively** to *predict the next token*.
- Trained on a **massive** amount of text data.



# Types of Transformers

- Encoder-Only Transformers:
  - Architecture: Consist only of the encoder stack.
  - Use Cases: Used for tasks where a fixed-length representation of the input is needed (e.g. text classification).
  - Example: BERT (Bidirectional Encoder Representations from Transformers).
- Decoder-Only Transformers:
  - Architecture: Consist only of the decoder stack.
  - Use Cases: Used for autoregressive tasks where the model generates output one token at a time.
  - Example: GPT (Generative Pre-trained Transformer).
- Encoder-Decoder Transformers:
  - Architecture: Both an encoder and a decoder.
  - Use Cases: Used for tasks that require sequence-to-sequence processing (e.g. machine translation).
  - Example: T5 (Text-to-Text Transfer Transformer).

# Zero- and Few-shot Contextual Learning

---

# The Fine-tuning Paradigm

## Fine-tuning

The model is trained via repeated gradient updates using a large corpus of example tasks.



# Zero-shot Contextual Learning

## Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.



# One-shot Contextual Learning

## One-shot

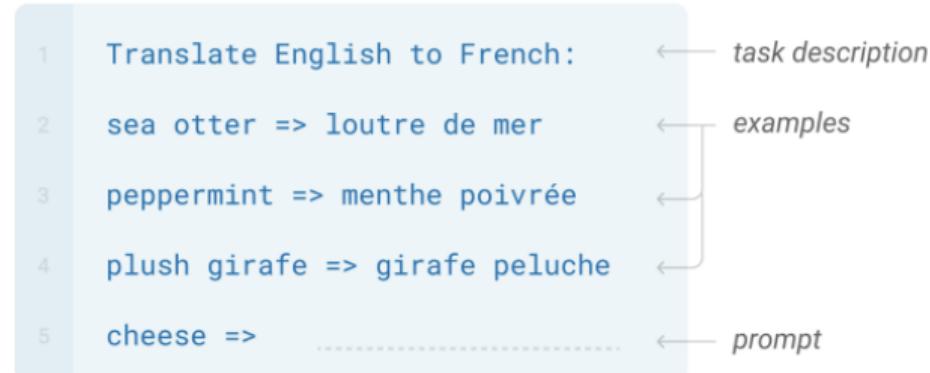
In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

- 1 Translate English to French: ← *task description*
- 2 sea otter => loutre de mer ← *example*
- 3 cheese => ← *prompt*

# Few-shot Contextual Learning

## Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.



# The GPT-3 Family

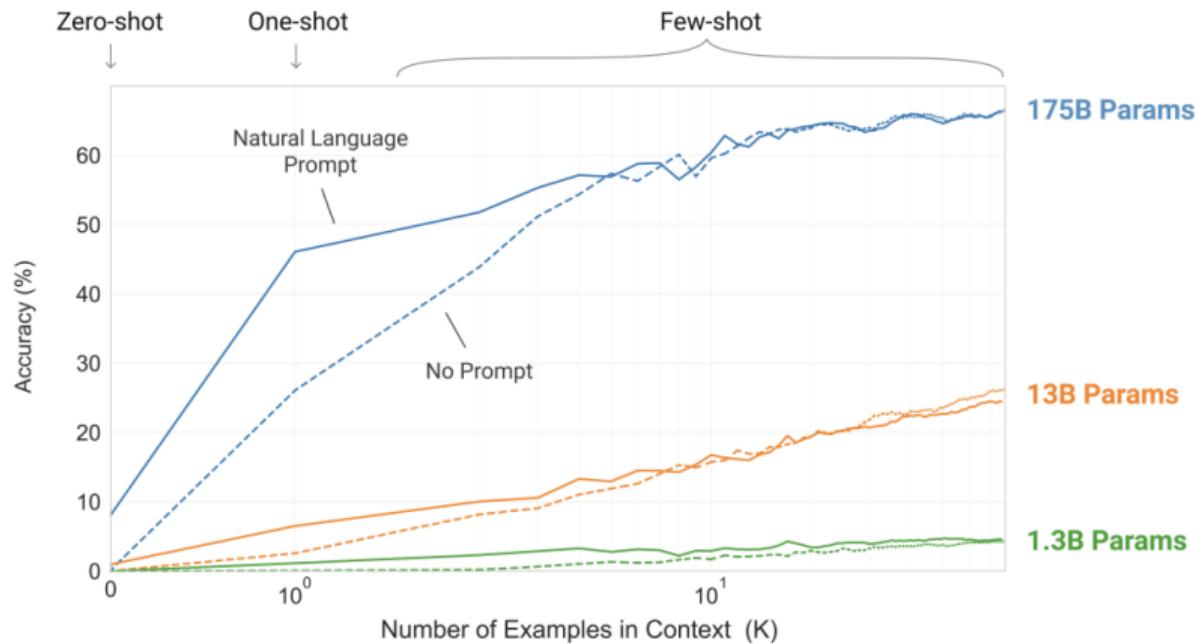
- How is this possible? **Scale** (in *parameters* and *training data*).

Dataset	Quantity (tokens)	Weight in training mix	Epochs elapsed when training for 300B tokens
Common Crawl (filtered)	410 billion	60%	0.44
WebText2	19 billion	22%	2.9
Books1	12 billion	8%	1.9
Books2	55 billion	8%	0.43
Wikipedia	3 billion	3%	3.4

Model Name	$n_{\text{params}}$	$n_{\text{layers}}$	$d_{\text{model}}$	$n_{\text{heads}}$	$d_{\text{head}}$	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	$6.0 \times 10^{-4}$
GPT-3 Medium	350M	24	1024	16	64	0.5M	$3.0 \times 10^{-4}$
GPT-3 Large	760M	24	1536	16	96	0.5M	$2.5 \times 10^{-4}$
GPT-3 XL	1.3B	24	2048	24	128	1M	$2.0 \times 10^{-4}$
GPT-3 2.7B	2.7B	32	2560	32	80	1M	$1.6 \times 10^{-4}$
GPT-3 6.7B	6.7B	32	4096	32	128	2M	$1.2 \times 10^{-4}$
GPT-3 13B	13.0B	40	5140	40	128	2M	$1.0 \times 10^{-4}$
GPT-3 175B or “GPT-3”	175.0B	96	12288	96	128	3.2M	$0.6 \times 10^{-4}$

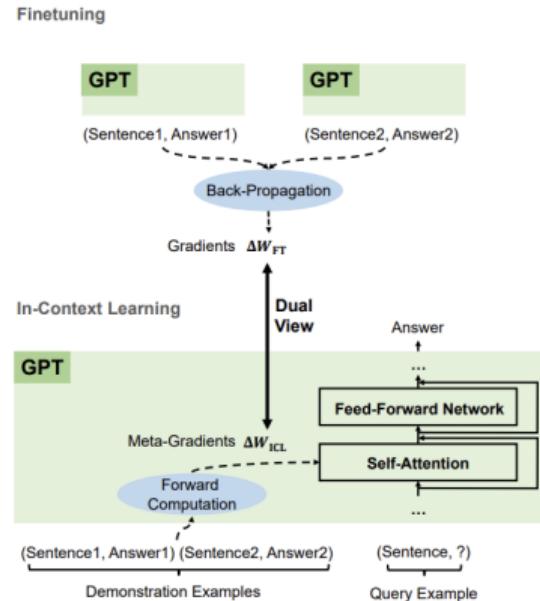
# GPT-3 and the Miracle of Scale

- Something **very special** happens at over a **hundred billion** parameters.



# WTF is going on here?

- How does GPT-3 (175B) do this without ever performing a gradient update?



Dai et al, "Why Can GPT Learn In-Context? Language Models Secretly Perform Gradient Descent as Meta-Optimizers."

GPT-X is Not ChatGPT

---

# ChatGPT and Conversational Agents

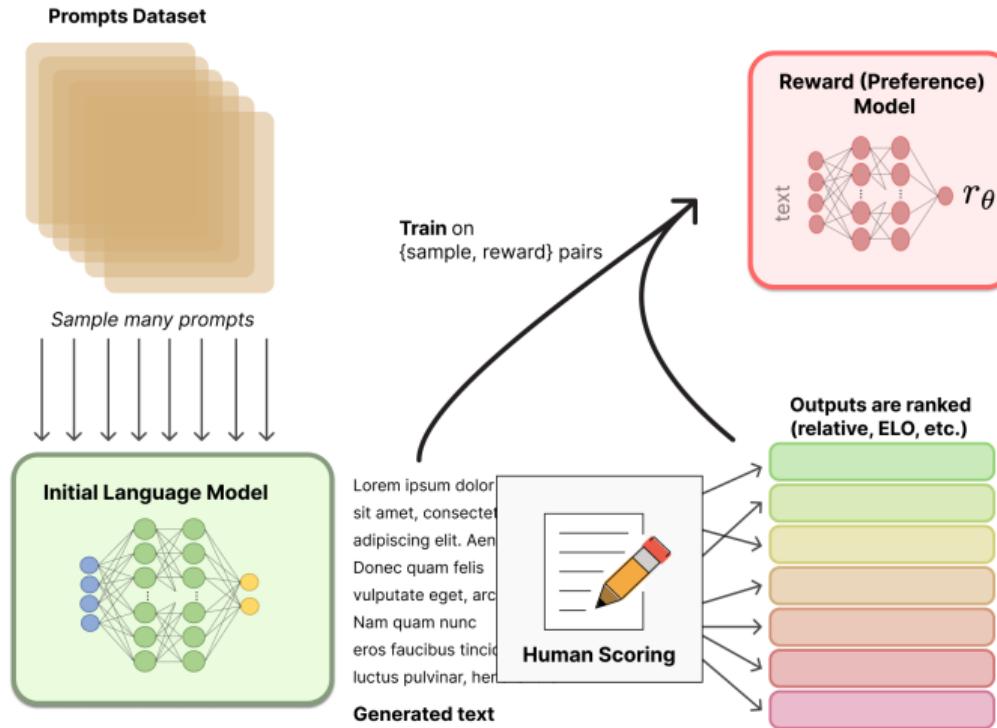
- GPT-3 is, in the end, a **very large** but still **autoregressive** model.
- It can generate **diverse** and **compelling** text from human prompts
- However, what makes a **good** text is hard to define and is subjective and context-dependent
- As the examples we have seen show, ChatGPT can:
  - Write stories that emulate styles and cite references (that may not exist!).
  - Produce **executable code snippets** that (probably?) do what the prompt asks.
  - (**Attempt to**) verify the **truthfulness** of what it produces.
- Writing a loss function to capture these characteristics is **intractable**, and our best language models are only trained for next token prediction.
- How does **ChatGPT** manage to do this?

# Reinforcement Learning to the Rescue

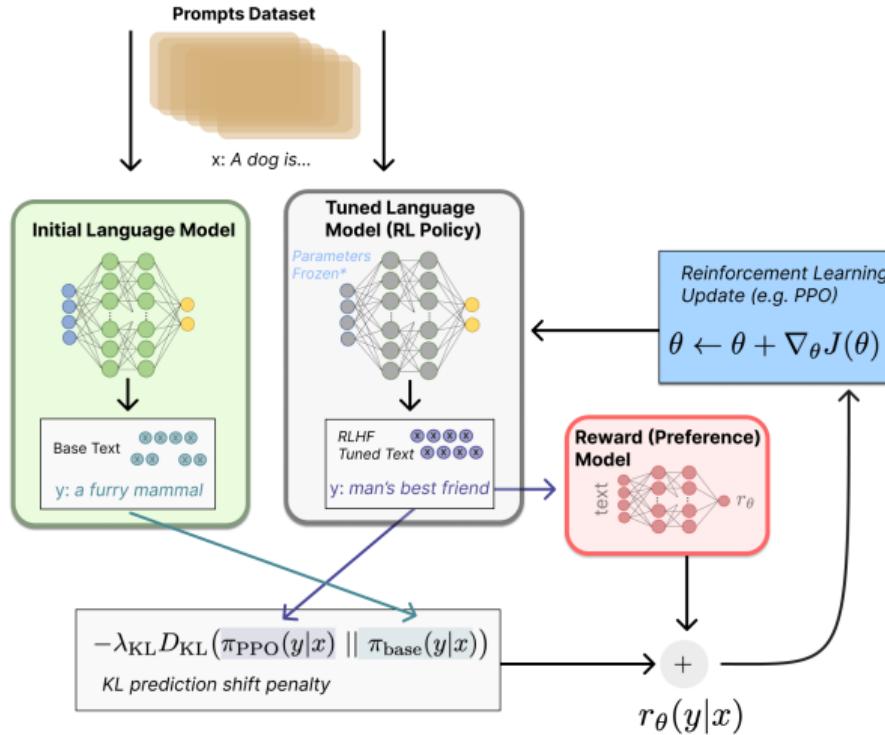
- OpenAI uses **reinforcement learning** to fine-tune their model to support high-quality interactions.
- They treat the LLM as an **agent** that:
  1. Observes the **state** – a **context** and possibly a **prompt**.
  2. Produces an **action** – the **answer** to the current contextualized prompt.
- It can then be trained with **any** reinforcement learning algorithm – they use Proximal Policy Optimization (PPO).
- This neatly sidesteps the need to define a supervised **loss function**.
- But... Don't we still need some sort of **reward**?

*Schulman et al, "Proximal Policy Optimization Algorithms."*

# Inverse Reinforcement Learning to the Rescue!



# Reinforcement Learning from Human Feedback



## Discussion

---

## Leftovers

- This course does not intend to be a **comprehensive** introduction to **all** of Machine Learning.
- There are **many** topics that could have been included, given **more time**.
- My hope is that the **fundamentals** you have acquired here are enough to allow you to **acquire** new skills and knowledge **on your own**.

## Leftover: Unsupervised Learning

- Some techniques from **unsupervised learning** are great to have in your toolbox:
  - **KMeans Clustering**: Allows you to **discover** structure in **unlabeled** data.
  - **Principal Component Analysis (PCA)**: Allows you to map **high-dimensional data** onto the principal axes of variation. Useful for **dimensionality reduction** and **noise removal**.
  - **Gaussian Mixture Models (GMMs)**: Allows you to **discover** structure in **unlabeled** data and explain it with a **statistical model**.

## Leftover: Recurrent Neural Networks (RNNs)

- What if our data has a **temporal dimension** of **variable size**?
- We already saw (very briefly) how **self-attention networks** (Transformers) can handle inputs of this type.
- Can we not apply **traditional** neural networks?
- The answer is **yes** (sort of):
  - A **recurrent** network is one that has **loops** (i.e. it is **not** a feed-forward network).
  - But wait, can we still use backpropagation in a graph that is **not** a DAG?
  - Yes: **Back Propagation through Time (BPTT)**

## Leftover: Generative Models, Few-shot Learning, Meta-learning, Continual Learning, ...

- Well, you get the idea...
- When all else fails:
  - What is the **input space**  $\mathcal{X}$ ?
  - What is the **output space**  $\mathcal{Y}$ ?
  - What is the **family of functions**  $\mathcal{H}$  that *makes sense* for my problem?
  - What is the **loss function**  $\mathcal{L}$  that also *makes sense*?
  - What **data**  $\mathcal{D}$  do I have to **learn** from?