



UNIVERSITÀ
DEGLI STUDI
FIRENZE



Parallel Programming

Prof. Marco Bertini



UNIVERSITÀ
DEGLI STUDI
FIRENZE



Python: multithreading



No-GIL threads and parallelism

From GIL to Free-Threaded Python

- Classic CPython uses the Global Interpreter Lock (GIL):
 - Only one thread executes Python bytecode at a time.
 - We have I/O-bound multithreading, doesn't allow CPU-bound work.
- PEP 703 introduced an optional free-threaded build where the GIL is disabled.
 - Python 3.13: experimental free-threaded builds.
 - Python 3.14: free-threaded build is fully supported and faster (specializing interpreter re-enabled).
 - Use uv to install it

Installing with uv

- Generic syntax:
 - uv python install 3.13t (3.13 free-threaded)
 - uv python install 3.14+freethreaded
- For a project:
 - uv python install 3.14+freethreaded
 - uv venv -p 3.14+freethreaded
 - source .venv/bin/activate # or
.venv\Scripts\activate on Windows

Ensuring GIL is off

Even in free-threaded Python, some C extensions can re-enable the GIL for safety.

- Force the GIL off at runtime:
 - Environment variable:
 - `export PYTHON_GIL=0`
 - Or command-line switch:
 - `python -X gil=0 your_script.py`
 - Or running using uv:
 - `uv run -p 3.14+freethreaded python -X gil=0 your_script.py`



CPU-bound vs. I/O-bound

- CPU-bound:
 - Runtime dominated by computations.
 - Faster CPU \Rightarrow faster program.
 - Examples: numerical kernels, cryptography, image processing.
- I/O-bound:
 - Runtime dominated by disk/network/DB access.
 - Faster disk/network \Rightarrow faster program.
- Before free-threaded Python:
 - Threads mainly useful for I/O-bound tasks (especially. Using Asyncio).
 - Multiprocessing required for CPU-bound parallelism.



Threads vs Processes in a No-GIL World

- Threads (within one process):
 - Cheap to create; share memory.
 - Excellent for fine-grained parallelism in a free-threaded interpreter.
- Processes:
 - Heavyweight; separate address spaces, separate interpreters.
 - Good for isolation & multi-process scaling.
- With free-threaded CPython 3.14:
 - Threads can now give true CPU parallelism, similar to C/Java, while using the familiar threading and concurrent.futures APIs.



UNIVERSITÀ
DEGLI STUDI
FIRENZE



Thread packages

Three Building Blocks

- threading: manual threads and synchronization primitives
- concurrent.futures: ThreadPoolExecutor and Future objects
- multiprocessing: process-based pools for parallelism





Cheat Sheet: Thread vs ThreadPool vs Pool



- `threading.Thread`:
 - Full control over threads and synchronization
 - Good for few long-lived threads and custom architectures
- `ThreadPoolExecutor`:
 - High-level API with futures, ideal for many similar tasks
- `multiprocessing.Pool`:
 - Process-based parallelism, good for isolation and classic CPython



threading Module: Thread Lifecycle

- Create threads with
`threading.Thread(target=..., args=...)`
- Start with `.start()`, wait with `.join()`
- Pattern: build list of threads, start all, join all
- In free-threaded Python, these threads can run
CPU work in parallel



threading: Shared Data and Locks

- Multiple threads may access shared mutable state
- Use `threading.Lock()` to protect critical sections
- `with lock:` ensures `acquire()/release()` with exception safety
- Without locks: data races and incorrect results in no-GIL Python



threading: Other Useful Primitives

- RLock: re-entrant lock for recursive acquisitions
- Semaphore: allow limited number of concurrent threads
- Event: simple flag with wait()/set()/clear()
- Condition and Barrier: coordination between groups of threads





Using Functions with Manual Threads

- Basic idea: each thread calls the same function with different args
- Use a small wrapper to store results in a shared list/dict
- Example pattern:
 - `def worker(name, start, end):` compute and store result
 - `threading.Thread(target=worker, args=(...))`
- Good when logic naturally lives in a free function



Example

```
import threading

def count_in_range(name: str, start: int,
end: int) -> int:
    total = 0
    for x in range(start, end):
        total += x
    print(f"{name}: done")
    return total

def main():
    results = {}

    def worker(name: str, s: int, e: int):
        results[name] =
            count_in_range(name, s, e)

    threads: list[threading.Thread] = []
        for i in range(4):
            s = i * 1_000
            e = (i + 1) * 1_000
            t =
                threading.Thread(target=worker,
                                  args=(f"t{i}", s, e))
            threads.append(t)
            t.start()

        for t in threads:
            t.join()

        print("Total:", sum(results.values()))

    if __name__ == "__main__":
        main()
```



Example

```
import threading

def count_in_range(name: str, start: int,
                   end: int) -> int:
    total = 0
    for x in range(start, end):
        total += x
    print(f"{name}: done")
    return total

def main():
    results = {}

    def worker(name: str, s: int, e: int):
        results[name] =
            count_in_range(name, s, e)

    threads: list[threading.Thread] = []
    for i in range(4):
        s = i * 1_000
        e = (i + 1) * 1_000
        t =
            threading.Thread(target=worker,
                             args=(f"t{i}", s, e))
        threads.append(t)
        t.start()

    for t in threads:
        t.join()

    print("Total:", sum(results.values()))
```

```
if __name__ == "__main__":
    main()
```

Pattern highlights:

threads: list[threading.Thread] Each thread runs a function with different arguments.

Pure functions are easy to run in threads (few shared variables).

Use a wrapper function (worker) to store results in a shared dict/list.



Using Classes with Manual Threads



- Encapsulate state + behavior in a class
- Each instance has its own configuration and result field
- Pass bound method as target:
`threading.Thread(target=instance.run)`
- Easier to extend with logging, config, error handling





Example

```
import threading

class RangeCounter:
    def __init__(self, name: str, start:
                 int, end: int):
        self.name = name
        self.start = start
        self.end = end
        self.result = 0

    def run(self) -> None:
        total = 0
        for x in range(self.start,
                        self.end):
            total += x
        self.result = total
        print(f"{self.name}: done")

def main():
    workers = [
        RangeCounter("t0", 0, 1_000),
```

```
        RangeCounter("t1", 1_000, 2_000),
    ]

    threads = [
        threading.Thread(target=w.run)
        # bound method as target
        for w in workers
    ]

    for t in threads:
        t.start()
    for t in threads:
        t.join()

    total = sum(w.result for w in
                workers)
    print("Total:", total)

if __name__ == "__main__":
    main()
```



Example

```
import threading

class RangeCounter:
    def __init__(self, name: str, start: int, end: int):
        self.name = name
        self.start = start
        self.end = end
        self.result = 0

    def run(self) -> None:
        total = 0
        for x in range(self.start, self.end):
            total += x
        self.result = total
        print(f"{self.name}: done")

def main():
    workers = [
        RangeCounter("t0", 0, 1_000),
        RangeCounter("t1", 1_000, 2_000),
    ]
    threads = [
        threading.Thread(target=w.run)
        # bound method as target
        for w in workers
    ]
    for t in threads:
        t.start()
    for t in threads:
        t.join()
    total = sum(w.result for w in workers)
    print("Total:", total)
```

Key ideas:

You want to **encapsulate** state + behavior in a class.
Use bound methods (e.g. w.run) as target for Thread.
State (start, end, result) lives inside each instance ⇒ less shared data.
Easy to add more behavior (logging, error handling, configuration).

High-Level Threading: concurrent.futures

- ThreadPoolExecutor: manages a pool of worker threads
- submit(fn, *args): schedule work, returns a Future
- map(fn, iterable): parallel map-style execution
- Futures provide a unified interface for waiting on results



Example

```
from concurrent.futures import
ThreadPoolExecutor
from primes import count_primes
def task(args: tuple[int, int])
-> int:
    start, end = args
    return count_primes(start, end)
def main():
    ranges = [(1, 75_000), (75_000,
                           150_000), (150_000, 225_000),
                           (225_000, 300_000)]
```

with
ThreadPoolExecutor(max_workers=4)
as ex:
 results = list(ex.map(task,
 ranges))
 print("Total primes:",
 sum(results))

```
        if __name__ == "__main__":
            main()
```

Key idea: “apply this function to many inputs”.
Prepare a list of independent tasks.
Use `executor.map(fn, iterable)` for a clean “parallel map”.
The pool automatically:

- creates worker threads,
- schedules tasks,
- reuses threads between tasks.

Futures: What They Represent

- Future = handle to a not-yet-computed result
- `fut.result()`: return value or raise captured exception
- `fut.done()`, `fut.exception()`,
`fut.add_done_callback(cb)`
- Use `as_completed(futures)` to process tasks as they finish

Thread Pools: Map-Style Parallelism

- ThreadPoolExecutor is ideal for 'apply this function to many inputs'
- Prepare a list of independent tasks (e.g. ranges)
- Use `executor.map(fn, iterable)` for simple parallel map
- Pool handles thread creation, scheduling, and reuse

Thread Pools: Futures & as_completed

- More control with submit + as_completed
- Submit tasks:

```
futures = [executor.submit(fn, arg)
for arg in items]
```
- Iterate as_completed(futures) to process tasks as they finish
- Handle per-task exceptions via fut.result() inside try/except



Example

```
from concurrent.futures import
ThreadPoolExecutor, as_completed
from primes import count_primes

def main():
    ranges = [(1, 75_000), (75_000,
                           150_000), (150_000, 225_000),
                           (225_000, 300_000)]

    with
        ThreadPoolExecutor(max_workers=4) as ex:
            futures = {
                ex.submit(count_primes,
                          start, end): (start, end)
                for (start, end) in ranges
            }

            total = 0

            for fut in
                as_completed(futures):
                    start, end = futures[fut]
                    try:
                        n_primes = fut.result()
                        print(f"{start}-{end}:
                               {n_primes}")
                        total += n_primes
                    except Exception as e:
                        print(f"Error in range
                               {start}-{end}: {e}")

            print("Total primes:", total)

if __name__ == "__main__":
    main()
```



Example

```
from concurrent.futures import
ThreadPoolExecutor, as_completed
from primes import count_primes

def main():
    ranges = [(1, 75_000), (75_000,
                           150_000), (150_000, 225_000),
                           (225_000, 300_000)]

    with
        ThreadPoolExecutor(max_workers=4) as ex:
            futures = {
                ex.submit(count_primes,
                          start, end): (start, end)
                for (start, end) in ranges
            }

    total = 0
    for fut in as_completed(futures):
        start, end = futures[fut]
        n_primes = fut.result()
        print(f"{start}-{end}:
                           {n_primes}")
        total += n_primes
    print("Total primes:", total)

if __name__ == "__main__":
    main()
```

For more control when using ThreadPool, use `submit + as_completed`

Benefits:

Process results as soon as each task finishes.

Handle exceptions per-task (inside `try: fut.result()`).

Same pattern works with ProcessPoolExecutor if you want processes instead of threads.



Process Pools vs Thread Pools (Usage Focus)

- ThreadPoolExecutor:
 - Shared memory, low overhead, great on free-threaded Python
 - Use when tasks share in-memory data or objects
 - Use for CPU-bound pure Python when libraries are thread-safe
 - Very low overhead to send arguments/results (no pickling).
- multiprocessing.Pool / ProcessPoolExecutor:
 - Separate processes, pickling overhead: arguments/results are pickled thus must be serializable.
 - Use for isolation or on classic CPython for CPU-bound work

Process-Based Parallelism: multiprocessing.Pool

- Uses multiple worker processes rather than threads
- Each process has its own interpreter and memory space
- map / starmap run functions across workers
- Adds overhead (startup + pickling) but bypasses GIL on classic CPython





Expected Performance

- Classic CPython (with GIL):
 - Sequential: baseline
 - Threads / ThreadPool: ~same speed as sequential for CPU-bound
 - Multiprocessing: speed-up at the cost of overhead
- Free-threaded Python 3.14:
 - Threads and thread pools: real multi-core speed-up
 - Multiprocessing: still parallel, sometimes similar or slower than threads



When Threads Beat Multiprocessing

- Lower overhead: no process startup or pickling
- Shared memory: easy to share large read-only data
- Lower memory footprint: single interpreter, many threads
- Great default for CPU-bound pure Python on free-threaded CPython



When Multiprocessing Still Wins

- Process isolation: worker crash doesn't kill main program
- Works on classic CPython where GIL blocks CPU-bound threads
- Useful for libraries not yet compatible with free-threaded Python
- Natural building block for multi-machine / distributed setups

Race Conditions & Synchronization (Recap)

- Removing the GIL exposes real data races on shared state
 - e.g.: unprotected updates (eg. counter $+= 1$) can lose increments
- Locks and other primitives are required for correctness
- Free-threaded Python behaves like C/Java with respect to races

Measuring Performance Correctly

- Use `time.perf_counter()` for wall-clock timing
- Run each scenario multiple times, consider median
- Avoid mixing heavy I/O with CPU-bound work in benchmarks
- For deeper analysis: `timeit`, `cProfile`, line profilers, IDE tools



Summary & Takeaways

- Python 3.14 free-threaded builds enable real CPU-bound multithreading
- threading + concurrent.futures give powerful, familiar APIs
- multiprocessing remains valuable for isolation and legacy setups
- Same code can run on classic and free-threaded Python, with different speedups

Books

- Python Parallel Programming Cookbook, G. Zaccone, Packt Publishing, Chapt. 2
- The Python 3 Standard Library by Example, D. Hellmann, Addison Wesley, Chapt. 10

