

FUNDAMENTALS OF MACHINE LEARNING

AA 2023-2024

Prova Finale (FACSIMILE)

18 Dicembre, 2023

Istruzioni: Niente libri, niente appunti, niente dispositivi elettronici, e niente carta per appunti. Usare matita o penna di qualsiasi colore. Usare lo spazio fornito per le risposte.

Instructions: No books, no notes, no electronic devices, and no scratch paper. Use pen or pencil. Use the space provided for your answers.

This exam has 5 questions, for a total of 100 points and 10 bonus points.

Nome: _____

Matricola: _____

1. **Multiple Choice:** Select the correct answer from the list of choices.

- (a) [5 points] True or False: A K-nearest neighbor classifier is only able to learn linear discriminant functions. ☐ True ☐ False
- (b) [5 points] True or False: Projecting a dataset onto its first principal component maximizes the variance of the projected data. ☐ True ☐ False
- (c) [5 points] True or False: The K-means algorithm is guaranteed to find the best cluster centers for any dataset. ☐ True ☐ False
- (d) [5 points] True or False: A Parzen kernel density estimator uses only the nearest sample in the dataset to estimate the probability of an input sample \mathbf{x} . ☐ True ☐ False
- (e) [5 points] How many parameters will a Multilayer Perceptron (MLP) for binary classification with a single hidden layer of width 10 and an input dimensionality of 8 have?
☐ 80 ☐ 99 ☐ 88 ☐ None of the above
- (f) [5 points] What will the entries of the Gram matrix be for a linear kernel?
 - ☐ $K[i, j] = (\mathbf{x}_i^T \mathbf{x}_j)^\gamma$
 - ☐ $K[i, j] = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|_2^2)$
 - ☐ $K[i, j] = \mathbf{x}_i^T \mathbf{x}_j$
 - ☐ None of the above
- (g) [5 points] Which of the following loss functions is called the negative log likelihood?
 - ☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C (\ln y_c - \ln \hat{y}_c)^2$
 - ☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C (y_c - \ln \hat{y}_c)^2$
 - ☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C y_c \ln \hat{y}_c$
 - ☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C \ln \hat{y}_c$
- (h) [5 points] How many iterations of gradient descent must we perform for an epoch of minibatch Stochastic Gradient Descent with a dataset of 1024 samples and a batch size of 16?
☐ 1024 ☐ 1 ☐ 32 ☐ 64

Total Question 1: 40

2. **Multiple Answer:** Select **ALL** correct choices: there may be more than one correct choice, but there is always at least one correct choice.

- (a) [5 points] What are the advantages of projecting data onto $K < D$ principal components?
- ☐ We eliminate noise in the original representation.
 - ☐ Classes are guaranteed to be linearly separable.
 - ☐ It is a nonlinear embedding that makes learning easy with simpler models.
 - ☐ Models trained on the reduced data are simpler.
- (b) [5 points] Which of the following are advantages of Ensemble Models (e.g. Committees)?
- ☐ They reduce the variance of the resulting model.
 - ☐ They are much more efficient than the base model.
 - ☐ They can reduce the expected error of the final model.
 - ☐ The resulting model is nonlinear even if the base model is linear.
- (c) [5 points] Which of the following are causes of the vanishing gradients when training neural networks?
- ☐ Saturated inputs to activation functions with near-zero derivatives when saturated.
 - ☐ Badly scaled input values.
 - ☐ Very deep models.
 - ☐ Bad random initialization of the network parameters.
- (d) [5 points] If we want to penalize classification errors less when training an SVM we should
- ☐ Increase the hyperparameter C .
 - ☐ Use a radial basis kernel.
 - ☐ Decrease the hyperparameter C .
 - ☐ None of the above.
- (e) [5 points] Which of the following are requirements for applying backpropagation to compute gradients in a deep network?
- ☐ The network must not be too deep.
 - ☐ The network must be a directed acyclic graph.
 - ☐ All activation functions must be differentiable.
 - ☐ All activation functions must be continuous.
- (f) [5 points] Which of the following are true of the Nadaraya-Watson estimator?
- ☐ It only requires some of the training data at test time.
 - ☐ It is a nonparametric method.
 - ☐ It estimates a nonlinear function of the input.
 - ☐ It estimates a linear function of the input.
- (g) [5 points] Which of the following models are nonparametric?
- ☐ The Multilayer Perceptron (MLP).
 - ☐ Logistic regression.
 - ☐ The K-Nearest Neighbor Classifier
 - ☐ Decision Trees.

Total Question 2: 35

3. [10 points] Show that a Committee Ensemble model using N bootstrapped linear regression models is a linear regression (i.e. that can be expressed as $\mathbf{w}^T \mathbf{x} + b$ for some \mathbf{w} and b).

4. [15 points] Show that a Multilayer Perceptron with two hidden layers with activation function $\sigma(x) = x$ is only capable of learning linear functions.

5. [10 points (bonus)] Design a Deep Convolutional Neural Network (with at least three convolutional layers and one or more pooling layers) to classify MNIST images (input size 28×28). Draw the network (or write pseudocode for its definition) and indicate how many parameters each layer has and the sizes of the intermediate feature maps.